

CSCI-620

Relational Databases

Roadmap

1. Overview of database engines
2. Data modeling
3. Entity-relationship modeling
4. Relational model
5. Modeling hints

Why use a database?

- ▶ Useful abstractions
- ▶ ACID
 - ▷ Atomicity
 - ▷ Consistency
 - ▷ Isolation
 - ▷ Durability

Database Engine

- ▶ Two main components
 - ▷ Storage manager
 - Authorization/integrity manager
 - Transaction manager
 - File manager
 - Buffer manager
 - ▷ Query processor

Storage Manager

- ▶ Interface between the DB and the OS
- ▶ Responsible for
 - ▷ Authorization
 - ▷ Interaction with the OS file system
 - Storage access
 - File organization
 - ▷ Efficient data storage/modification
 - Indexing and hashing
 - Buffer management

Transaction Manager

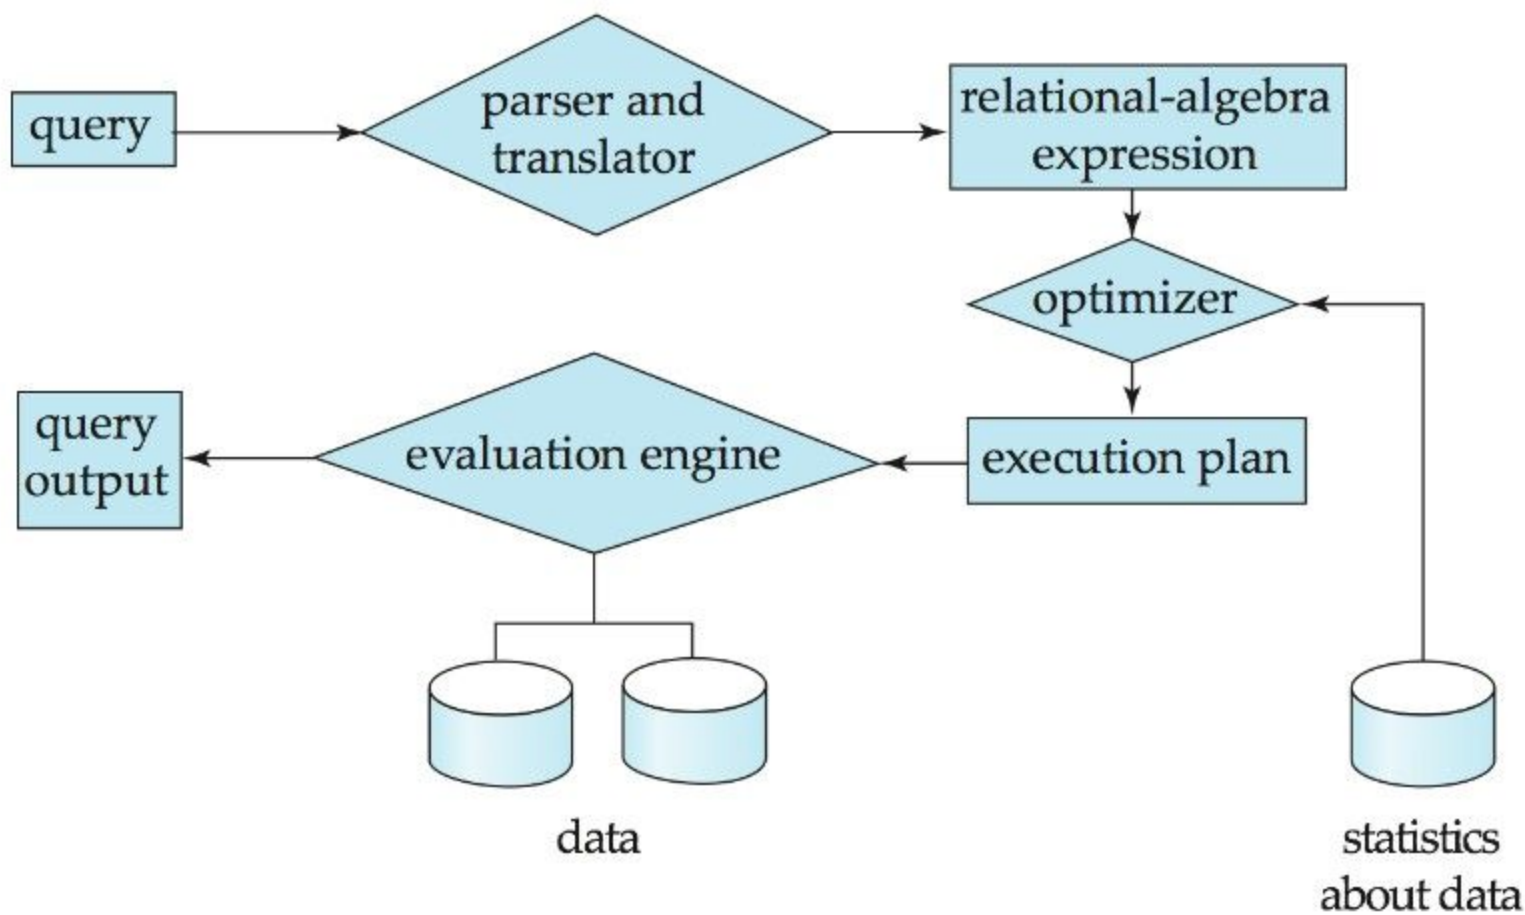
- ▶ Transaction management
 - ▷ Ensures the database is consistent if a failure occurs
 - ▷ “All or nothing” (atomic)
- ▶ Concurrency control
 - ▷ Makes sure multiple operations result in a consistent database

Transaction Manager

- ▶ Transactions are collections of operations for a single task
- ▶ Example
 - ▷ Assume a constraint balance > 0
 - ▷ Deduct 50 from A
 - ▷ Add 50 to the balance of B
 - ▷ Store the new balance

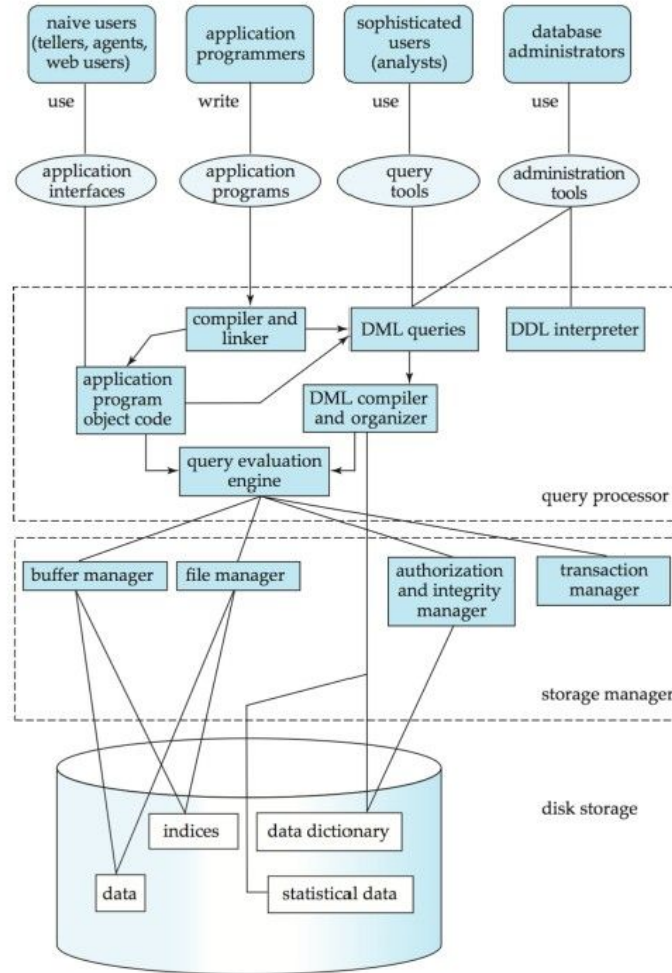
Query Processor

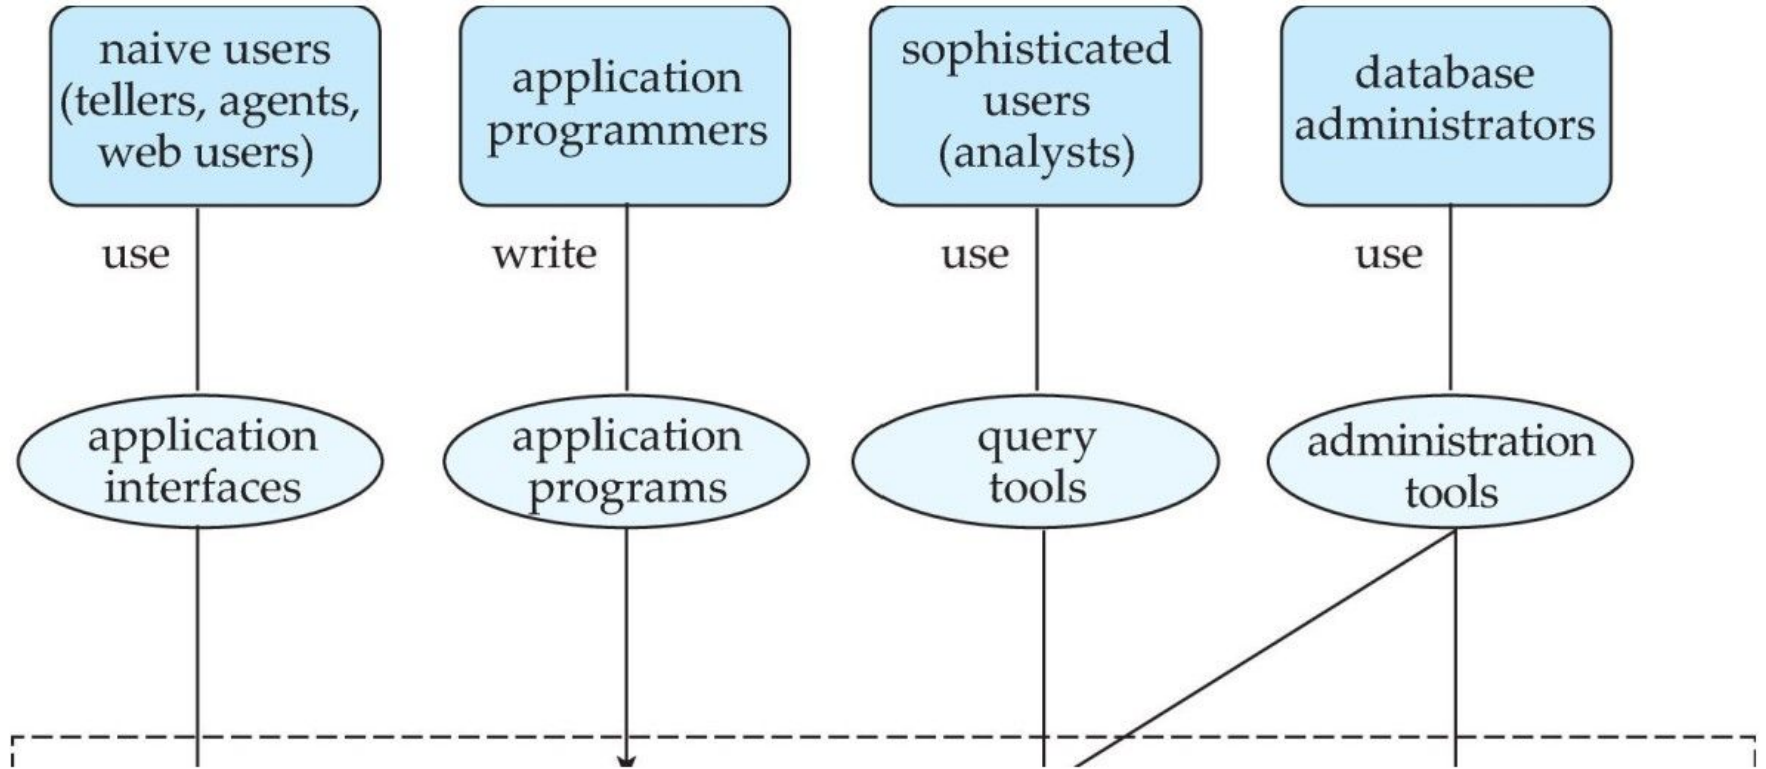
- ▶ Three major jobs
 - ▷ Parsing and translation
 - ▷ Optimization
 - ▷ Evaluation



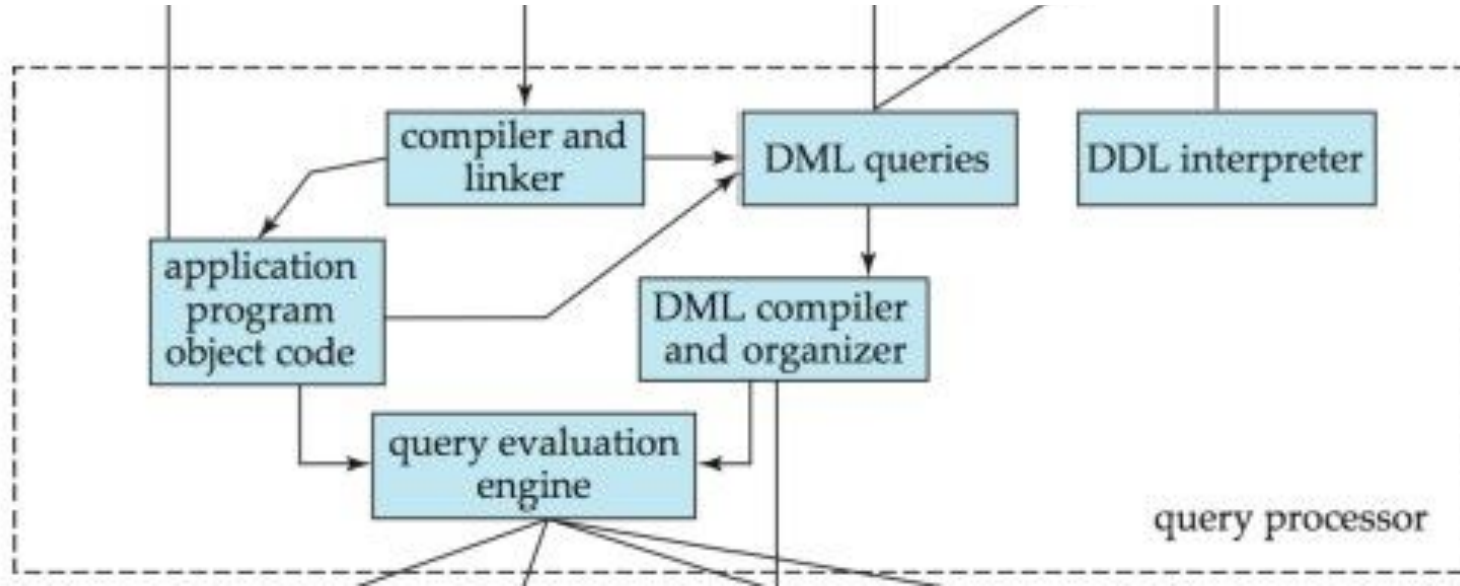
Database Engine

- ▶ All components work together
 - ▷ Storage manager must make sure transactions are durable
 - ▷ Query processor uses indexes managed by the storage manager
 - ▷ Transaction manager must provide consistent data to query processor

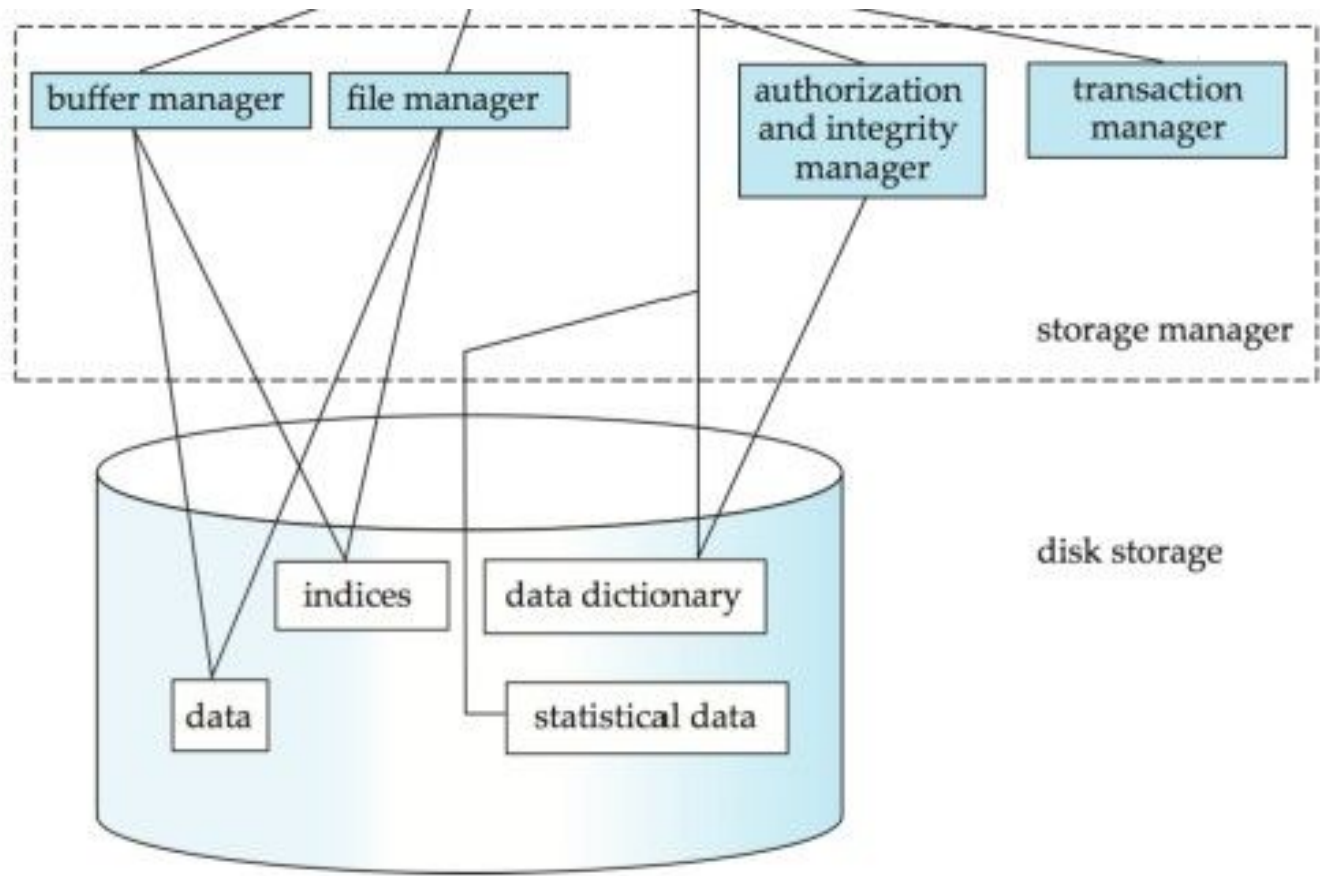




Interface



Query Processor



Storage Manager

SQL user
interface

forms
interface

report
generation
tools

data mining
and analysis
tools

front end

interface
(SQL API)

SQL engine

back end

Roadmap

1. Overview of database engines
2. **Data modeling**
3. Entity-relationship modeling
4. Relational model
5. Modeling hints

Levels of abstraction

1. View Layer

How applications access data
(hiding record details, more convenience, etc.)

2. Logical Layer

How data is stored in the database
(types of records, relationships, etc.)

3. Physical Layer

How data is stored on hardware
(actual bytes, files on disk, etc.)

Data Model

- ▶ A collection of tools for describing
 - ▷ Data
 - ▷ Data relationships
 - ▷ Data semantics
 - ▷ Data constraints

- ▶ Examples
 - ▷ Relational model
 - ▷ Entity relationship model
 - ▷ Object-based model
 - ▷ Semi-structured models (e.g. XML)
 - ▷ Network

Database Design

- ▶ What makes a *good* design?
- ▶ Conceptual modeling (e.g. ER model)
- ▶ Normalization theory
 - ▷ How do we check for bad designs?
 - ▷ How do we fix them?

Conceptual Model

- ▶ Before looking at the relational model, we need to have a way to think about what our database needs to store
- ▶ Many *conceptual* models exist that are independent of how a particular database stores data
- ▶ A common choice is the ER model

Conceptual Model

- ▶ Does not specify how data will actually be stored
- ▶ Also does not specify the interface we will use to access the data
- ▶ Useful for collecting requirements

Database Design


- ▶ Steps
 - ▷ Analyze requirements
 - ▷ Create a conceptual model
 - ▷ Create a logical model
 - ▷ Create a physical model

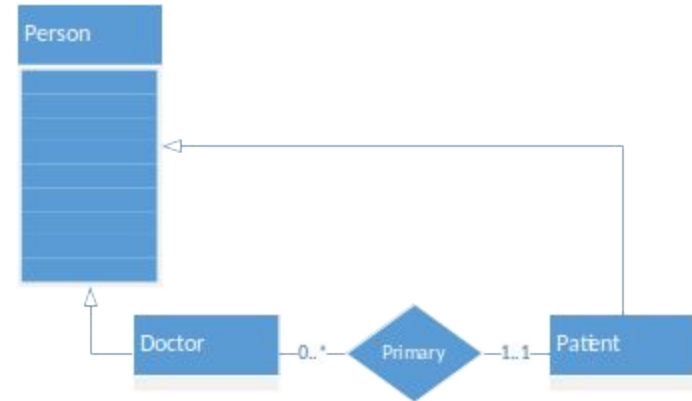
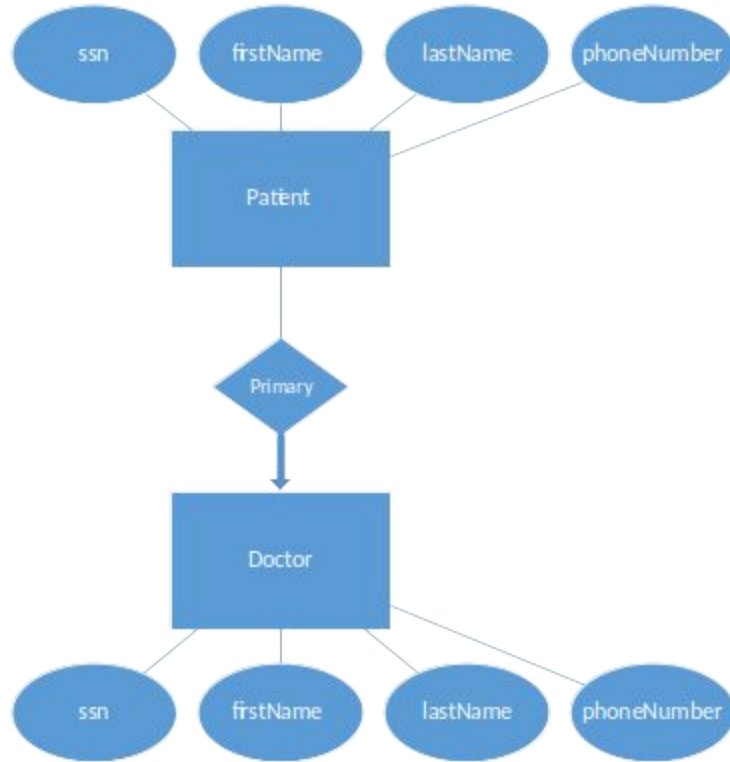
Roadmap

1. Overview of database engines
2. Data modeling
3. Entity-relationship modeling
4. Relational model
5. Modeling hints

Warning

- ▶ There are many versions of the ER model we will study
- ▶ We'll use the version from the textbook





Notation



Entity



Entity set

Entity Sets

Entity Sets

- ▶ A type of thing in the real world
- ▶ Distinguishable from other types
- ▶ Has a set of properties or attributes possessed by things of the same type

Singular



Shoe



Ball



Car

Plural



Shoes



Balls



Cars

Singular



Clock



Spoon



Cup

Plural



Clocks



Spoons



Cups

Single or plural names?



Person	
ssn	
firstName	
middleName	
lastName	
phoneNumber	
birthDate	
gender	
email	
occupation	

Attributes

Attributes

- ▶ Each attribute has an associated type which is normally *atomic* (indivisible)
- ▶ One or more attributes called the *primary key* uniquely identify an entity
- ▶ The set of valid values for an attribute is called the *domain*
- ▶ Attributes may be *simple* or *complex*

Person

ssn

firstName

middleName

lastName

phoneNumber

birthDate


gender

email

occupation



Relationship



Relationship set

Relationships

**Let's try building
a model**



Doctor

Clinical trial

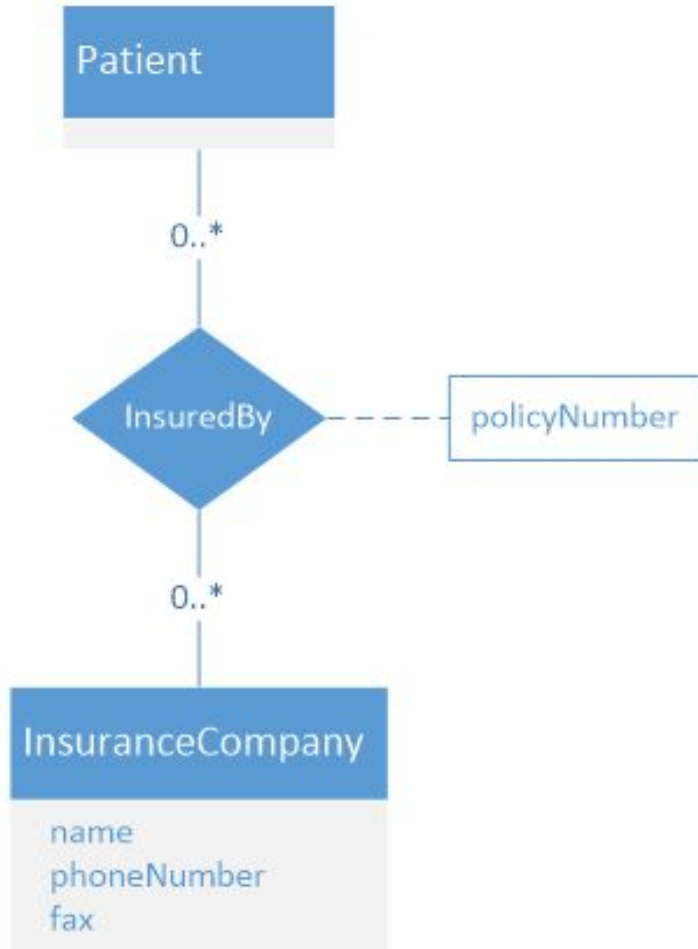


Patient

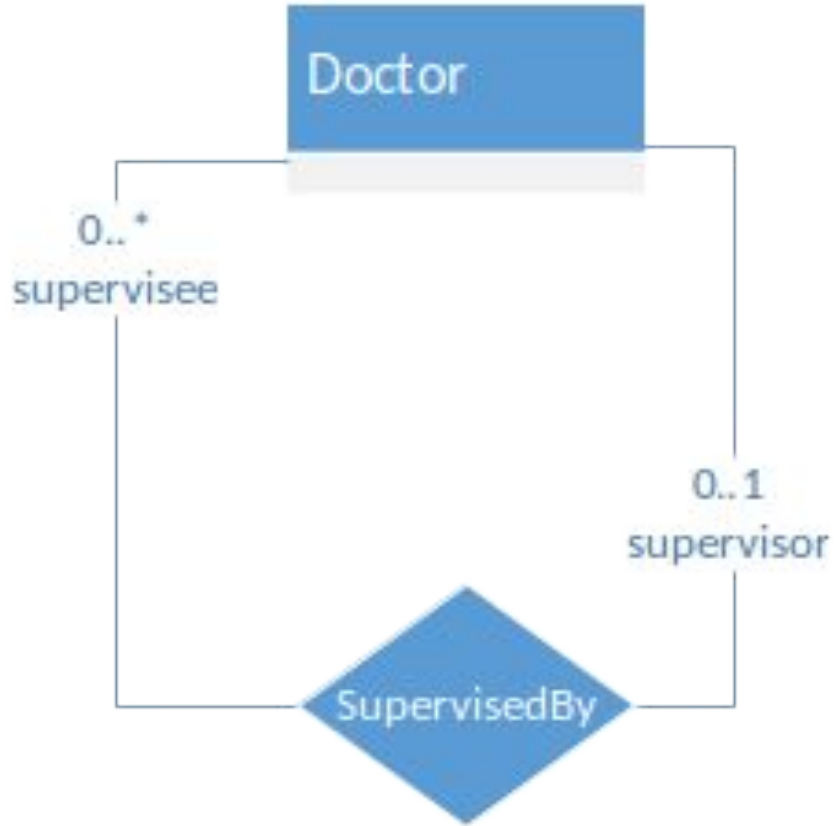
Non-binary Relationships

Relationships


- ▶ May also have attributes
- ▶ Each relationship has a *cardinality* or a restriction on the number of entities
- ▶ Implicitly defines a *role* for each entity set in the relationship




Relationship Set Attributes



Roles



1 to 1




1 to N




M to N

Cardinalities




One to one




(a)

Many to one

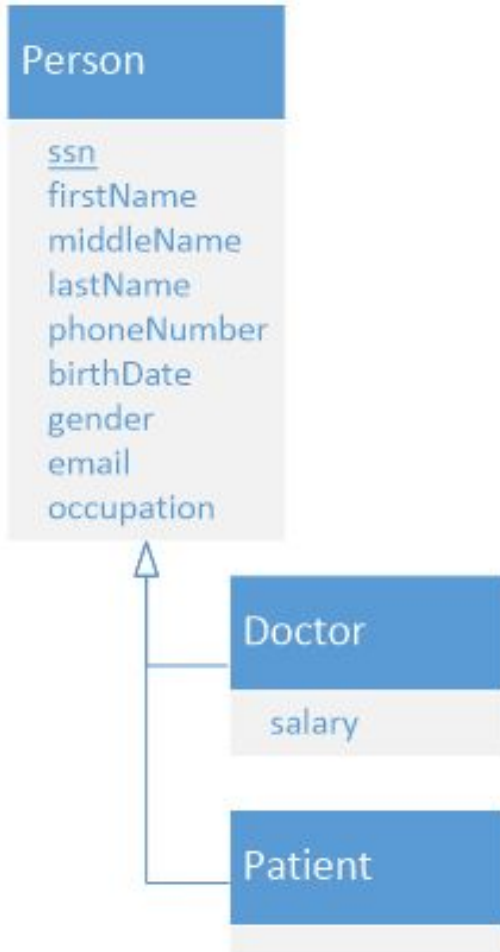


(b)

One to many



Many to many



Inheritance

Identity

- ▶ Each entity in an entity set must have some type of key
- ▶ This is usually a subset of the attributes associated with an entity
- ▶ The key should (never) or rarely change



Super



Candidate



Primary

Types of keys

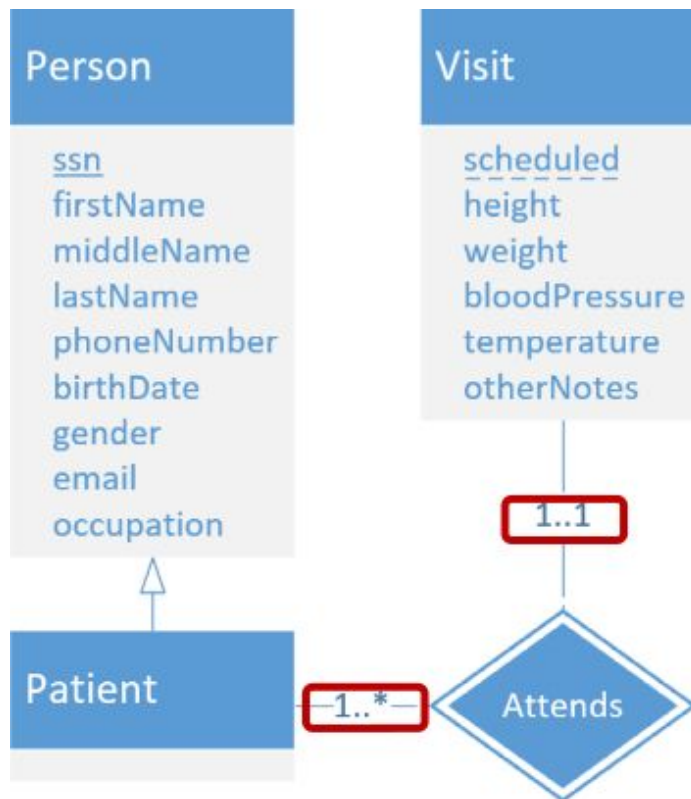
**What if there is no
candidate key?**

Solution 1

- ▶ Create an artificial ID attribute
- ▶ Ensure that a unique value is assigned
- ▶ Examples
 - ▷ Order number
 - ▷ Customer ID


Solution 2

- Use a *weak entity set*






VS.



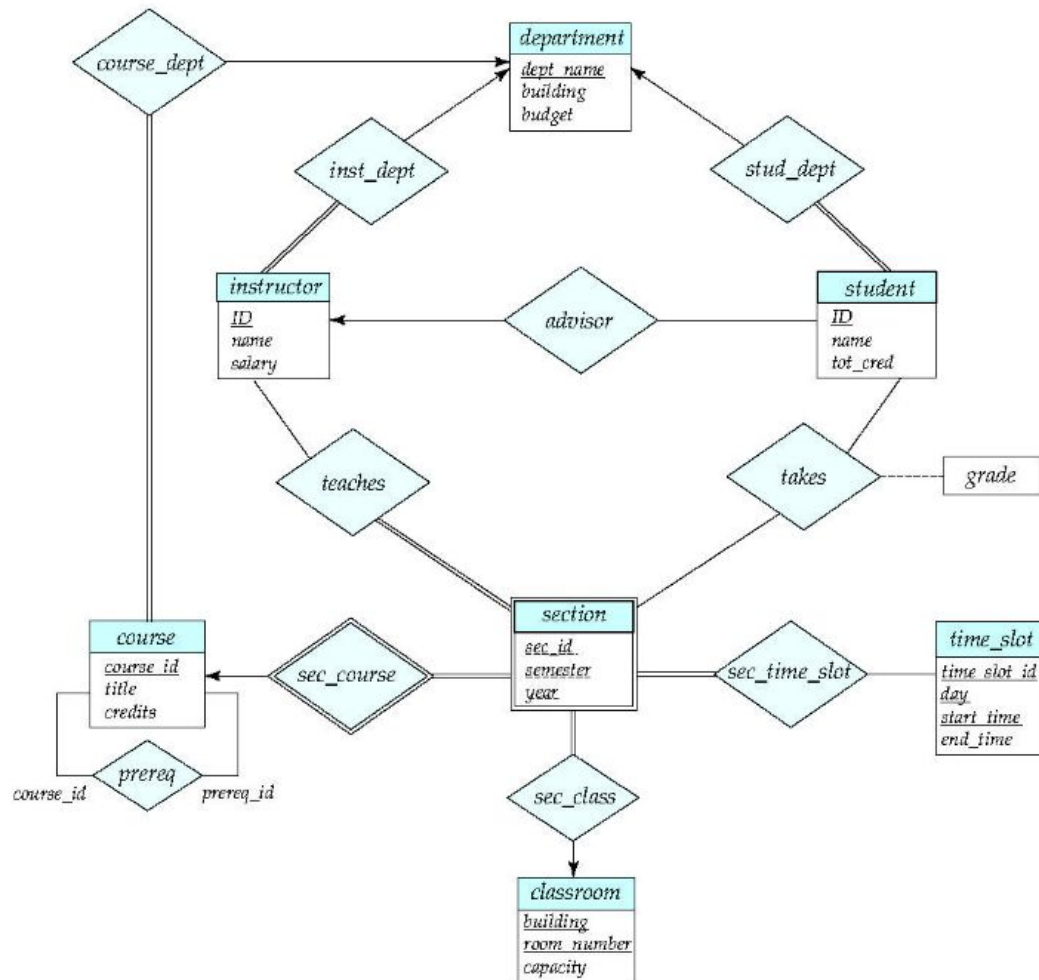
Design Decisions



Design Decisions

Design Decisions

- ▶ Attributes or entity sets?
- ▶ Entity sets or relationships?
- ▶ Strong or weak entity sets?
- ▶ Specialization/generalization



Reminder

- ▶ Remember that the ER model is *conceptual* and not what a database actually uses
- ▶ We need a more concrete model to actually implement our application
- ▶ Coming up next, the relational model!




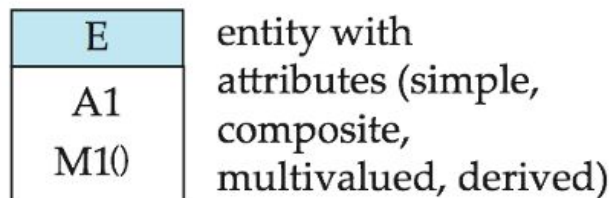
**Let's try more
modeling!**

UML


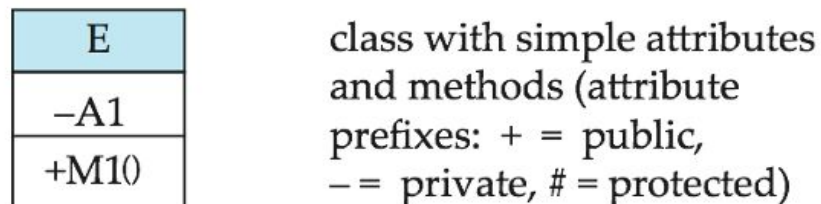
- ▶ Unified Modeling Language
- ▶ Used to graphically model software systems (*could* replace ER)
- ▶ We will use ER for databases, but UML to model your project software






ER Diagram Notation






Equivalent in UML



ER Diagram Notation



Equivalent in UML



Roadmap

1. Overview of database engines
2. Data modeling
3. Entity-relationship modeling
4. **Relational model**
5. Modeling hints

What now?

- ▶ Database systems don't use the ER model directly (it's a *conceptual* model)
- ▶ We need to move on to the next stage and pick a *logical* model
- ▶ The first model we'll explore is the *relational* model

Object 1: Maintenance Report

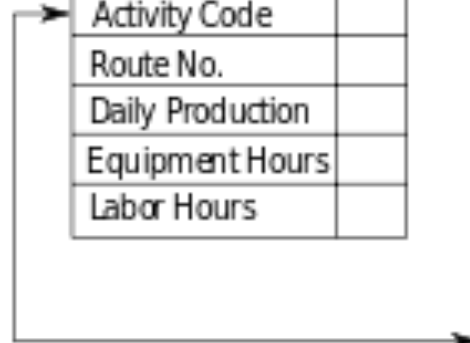
Date	
Activity Code	
Route No.	
Daily Production	
Equipment Hours	
Labor Hours	

Object 1 Instance

01-12-01
24
I-95
2.5
6.0
6.0

Object 2: Maintenance Activity

Activity Code	
Activity Name	
Production Unit	
Average Daily Production Rate	



Object-Relational Models

```
<Steps>
  <UserStep Code="Create" Start="NR" NotAssignedBehaviour="ERR" RequestComment="true"
    Help="true" Delegate="true" Iteration="1" CanSendToPool="true" StepAssignmentType="S">
    <Directions>
      <Direction Code="Direction1" To="Step1" Direction="F">
        <Condition Type="BS" Script="GT(WFDATA( "MAIN", "AMOUNT" ),"1000")" />
        <From>Create</From>
      </Direction>
      <Direction Code=" Direction2" To="Step2" Direction="F">
        <Condition Type="BS" Script="GT(WFDATA( "MAIN", " AMOUNT " ),"5000")" />
        <From>Create</From>
      </Direction>
      <Direction Code=" Direction3" To="Step3" Direction="F">
        <Condition Type="A" />
        <From>Create</From>
      </Direction>
    </Directions>
  </UserStep>
</Steps>
```

Semi-Structured Models

Patient				
ssn	firstName	middleName	lastName	primaryDoctor
235-14-7854	Sandra	null	Smith	943-23-9874
192-48-0924	John	Richard	Moore	862-74-3611

Doctor		
ssn	firstName	lastName
943-23-9874	Rachel	Wang
862-74-3611	Chris	Patel

Relation instances



Great idea!

Information Retrieval

A Relational Model of Data for Large Shared Data Banks

E. F. Codd
IBM Research Corporation, Yorktown Heights, New York

Large data banks must be designed from the beginning to know how the data is organized in the machine (the internal representation) by a prompting service which supplies such information is not a satisfactory solution. Activities of users at terminals and most program programs should remain unaffected when the internal representation of data is changed and even when some aspects of the external representation are changed. Changes in data representation will often be needed as a result of changes in query, update, and report traffic and natural growth in the types of stored information.

Existing noninferential, formatted data systems provide users with tree-structured files or slightly more general networks

REJECTED

At first...

Comments

- ▶ “... at first sight I doubt that anything complex enough to be of practical interest can be modeled using relations.”
- ▶ “... any realistic model might end up requiring dozens of interconnected tables—hardly a practical solution given that, probably, we can represent the same model using two or three properly formatted files.”
- ▶ “The paper can be safely rejected.”

However

- ▶ “... no real-world example (...) any model of practical interest can be cast in it.”
- ▶ “... no experiments (...) how it compares with traditional ones on real-world problems.”
- ▶ “... to extract any significant answer from any real database, the user will end up with the very inefficient solution of doing a large number of joins.”

Relations


- ▶ A relation is a group of related attributes like in an entity set
- ▶ Each relation should have a *primary key*
- ▶ Relations may also have *foreign keys* or attributes which refer to other relations
- ▶ In a relational database, these are represented as *tables*

Patient			
SSN	First Name	Middle Name	Last Name
235-14-7854	Sandra		Smith
192-48-0924	John	Richard	Moore
821-13-2108	Laura		Turner

Visit		
SSN	Scheduled	Weight
235-14-7854	09/03/2015	141.5
821-13-2108	10/18/2015	167.8



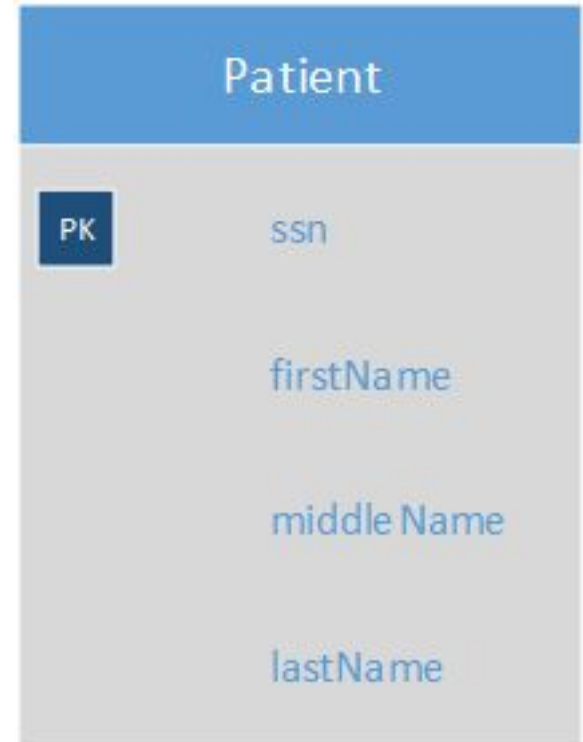
Patient-Visit					
SSN	First Name	Middle Name	Last Name	Scheduled	Weight
235-14-7854	Sandra		Smith	09/03/2015	141.5
821-13-2108	Laura		Turner	10/18/2015	167.8



Relational model usage

Strong Entity Sets


- ▶ Attributes of the entity set become attributes of the relation
- ▶ Primary keys are maintained



Patient				
ssn	firstName	middleName	lastName	primaryDoctor
235-14-7854	Sandra	null	Smith	943-23-9874
192-48-0924	John	Richard	Moore	862-74-3611

Doctor		
ssn	firstName	lastName
943-23-9874	Rachel	Wang
862-74-3611	Chris	Patel

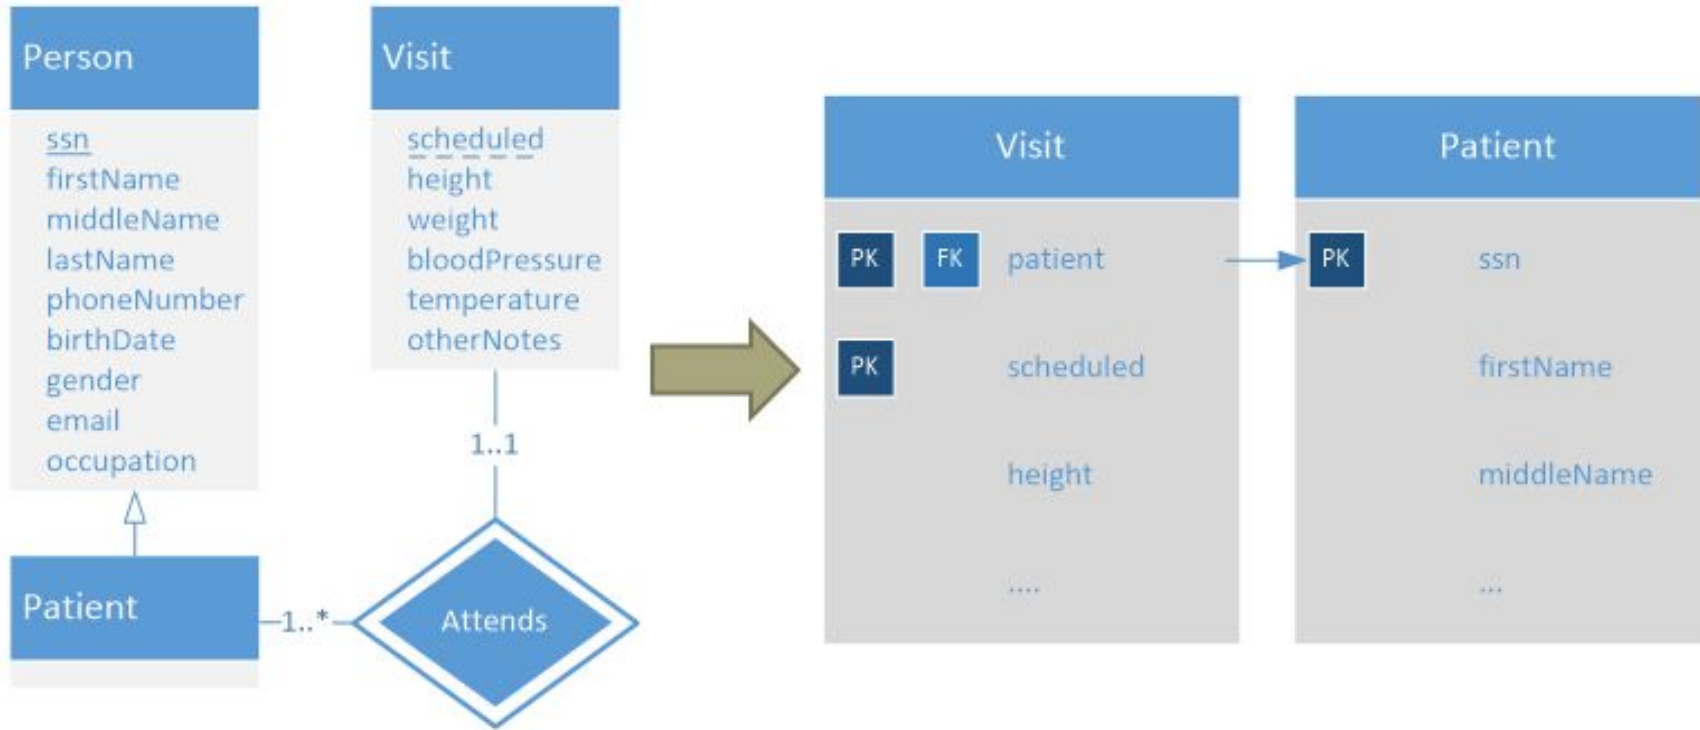
Relation Instances



Types

Types

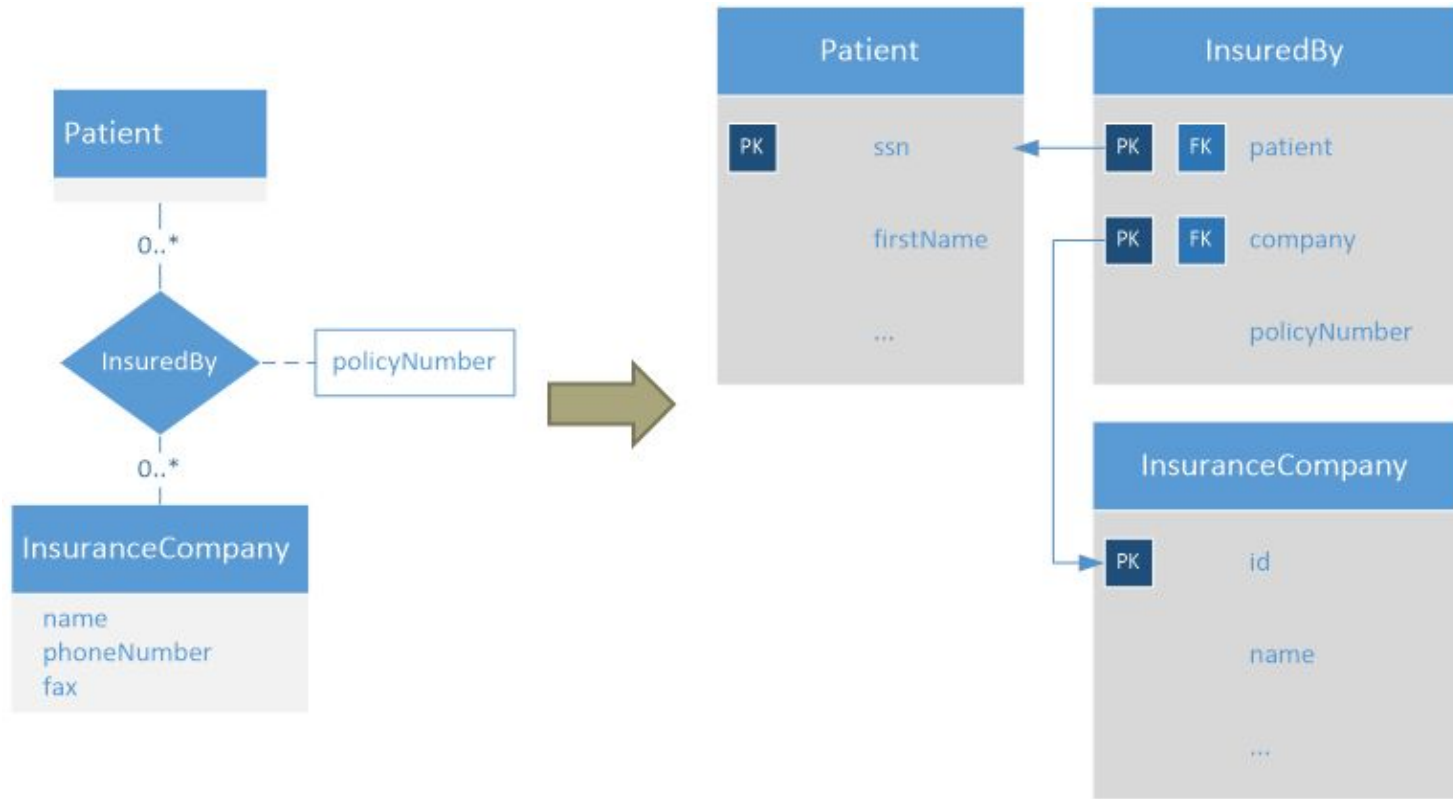
- ▶ String
 - ▷ Fixed (CHAR)
 - ▷ Variable length (VARCHAR, TEXT)
- ▶ Numbers
 - ▷ Integer (INTEGER)
 - ▷ Floating point (DECIMAL)
- ▶ Dates
 - ▷ Date (DATE)
 - ▷ Times (TIME)
 - ▷ Date and time (TIMESTAMP)



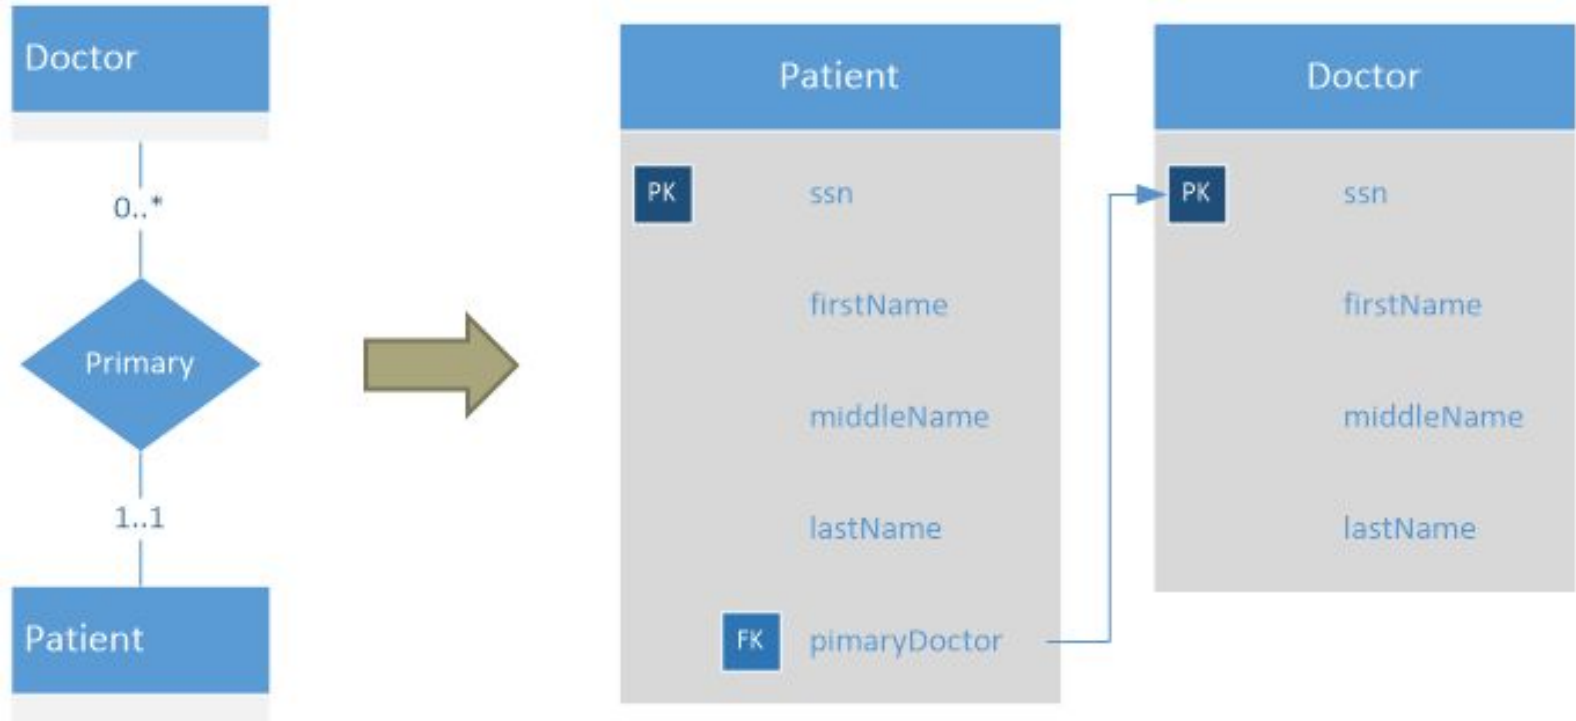
Weak Entity Sets

Weak Relationship Sets

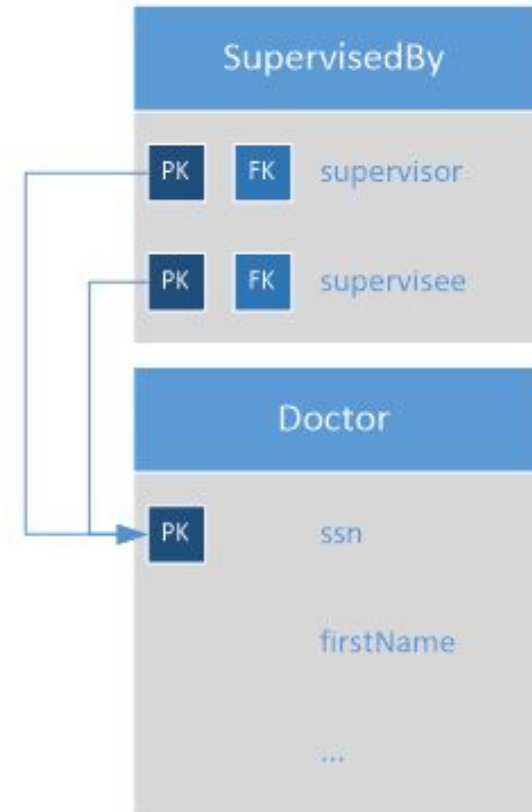

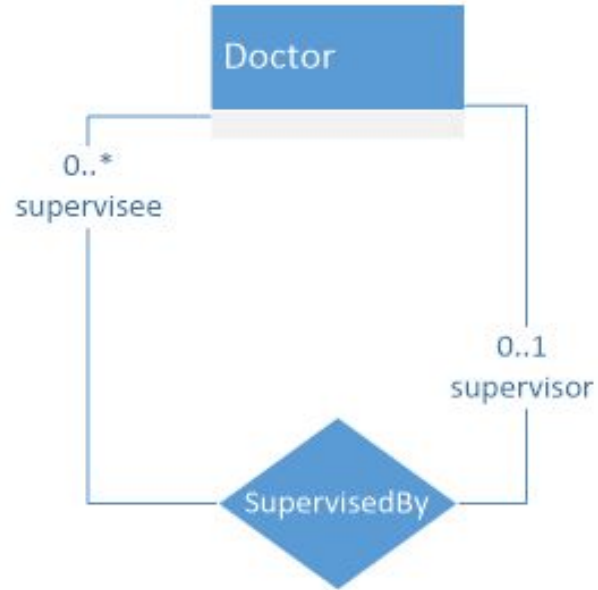
- ▶ We can discard these since the relationship is captured by the weak entity set
- ▶ Example: The **Attends** relationship is captured by the **Visit** relation created from the weak entity set **Visit**



Strong Relationships



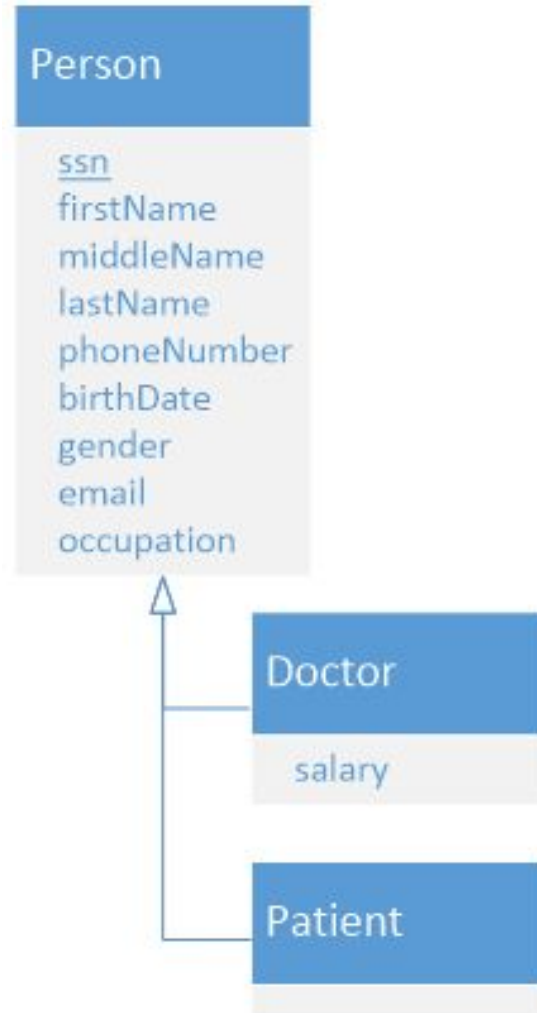
Optimization

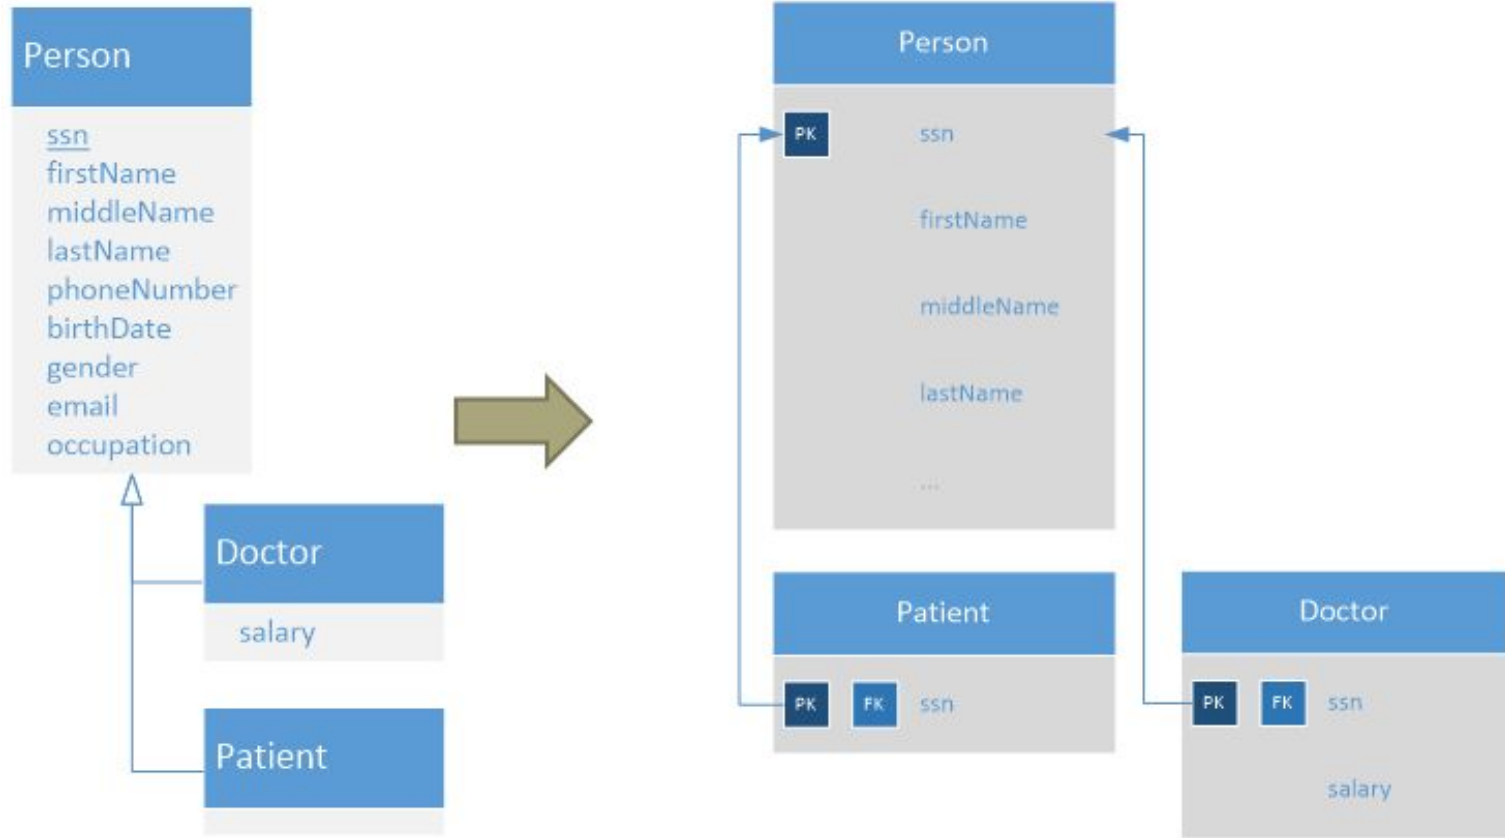


Roles

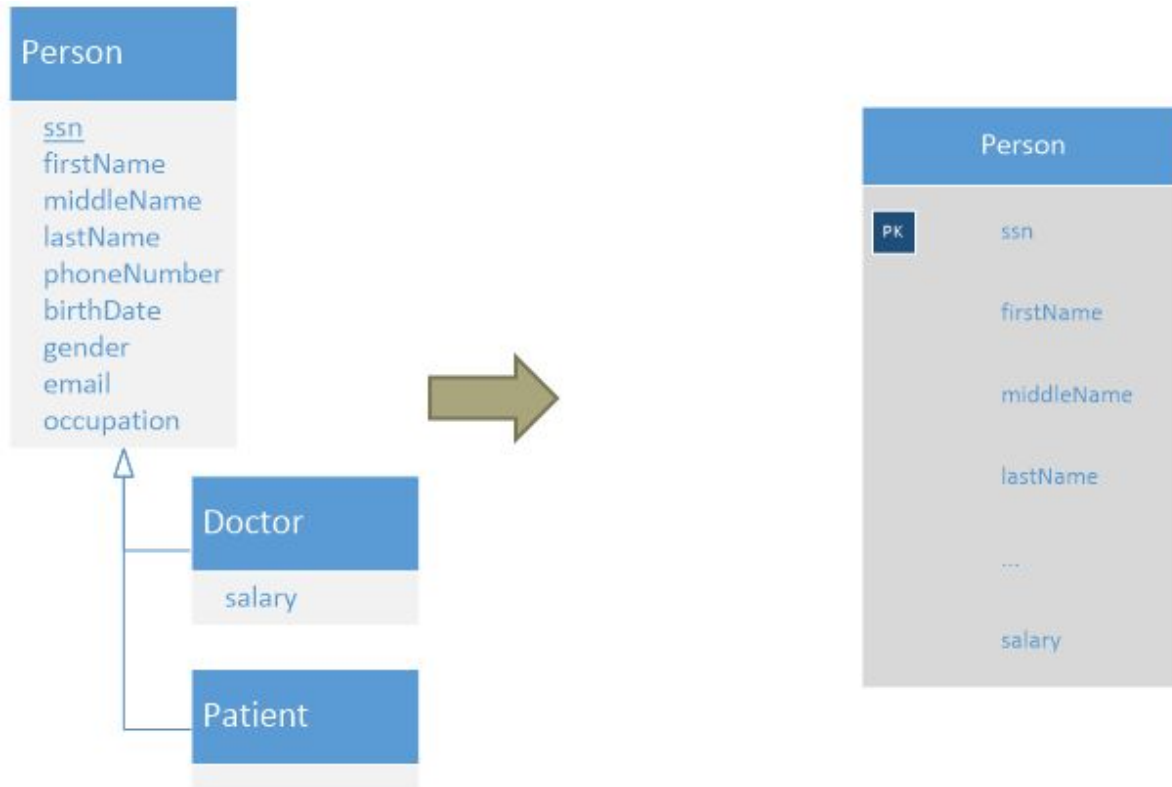
Inheritance Strategies

- ▶ Whole
- ▶ Top
- ▶ Bottom

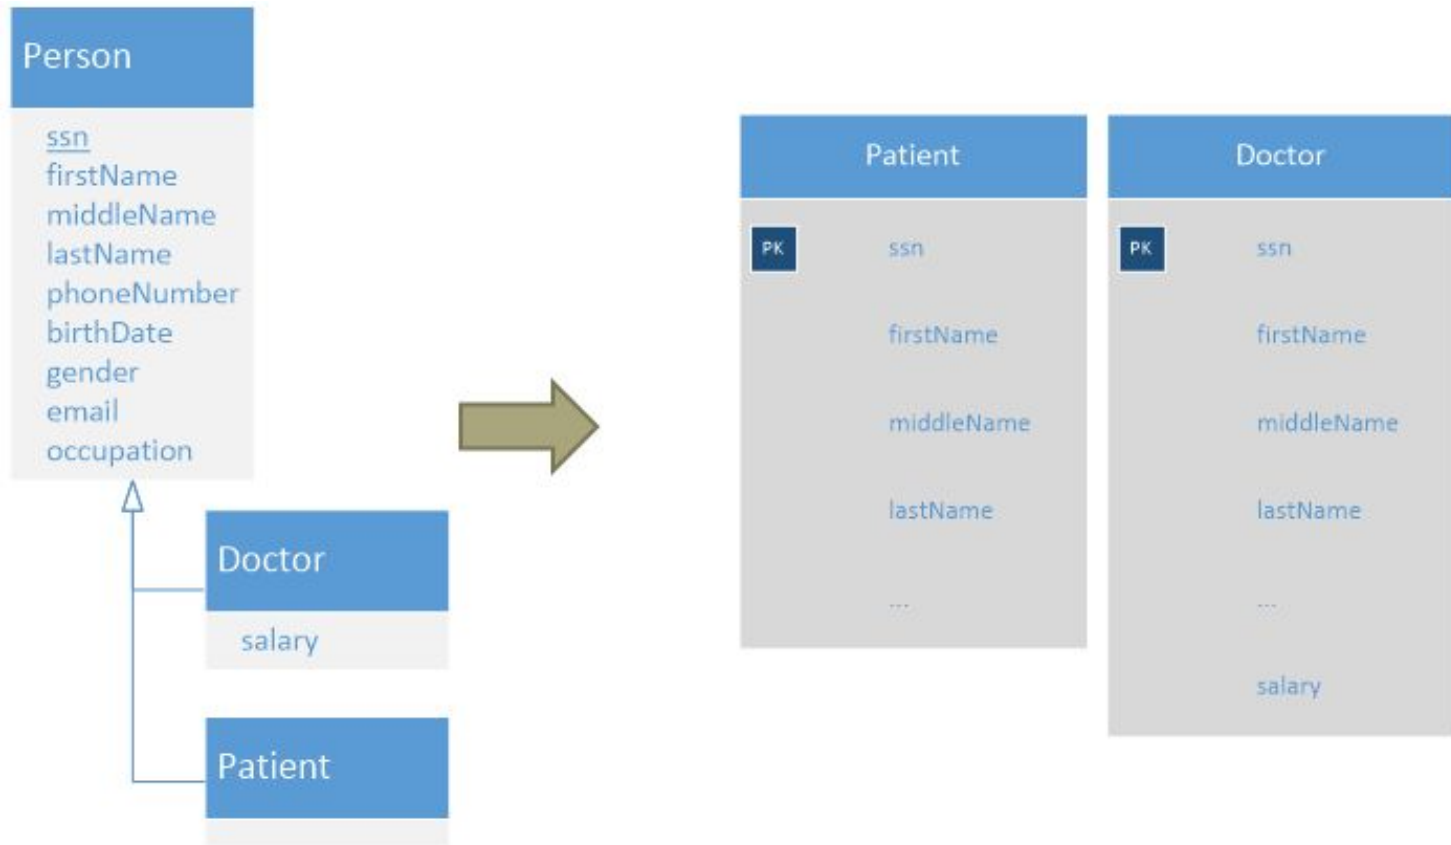




Whole hierarchy



Top of the hierarchy



Bottom of the hierarchy

Roadmap

1. Overview of database engines
2. Data modeling
3. Entity-relationship modeling
4. Relational model
5. Modeling hints



Use your brains!


Database Design

Bad smells in design




Database Design

- ▶ Hint #1: redundancy



Database Design

- ▶ Hint #2: incompleteness



- ▶ Hint #3: multiple models

Database Design



Drugs



Patients



Furniture

DDL

- ▶ Data Definition Language (DDL)
- ▶ Specific notation for defining database schema
- ▶ Similar to the Java compiler changing your class into a template

DDL

- ▶ Dictionary (metadata - data about data)
 - ▷ Database schema (logical)
 - ▷ Integrity constraints
 - Primary keys
 - Foreign keys
 - Constraints (e.g. $\text{balance} \geq 0$)
 - ...
 - ▷ Authorization - who can access what

Creating Tables

```
CREATE TABLE Patient (  
    ssn CHARACTER(9),  
    firstName VARCHAR(75) NOT NULL,  
    middleName VARCHAR(75),  
    lastName VARCHAR(75) NOT NULL,  
    primaryDoctor CHARACTER(9),  
    PRIMARY KEY (ssn),  
    FOREIGN KEY (primaryDoctor)  
    REFERENCES Doctor(ssn)  
);
```