

CSCI-620

NoSQL

Roadmap

1. Intro to NoSQL
2. Document databases
3. Key-value databases
4. Wide column stores
5. Graph databases



NoSQL is not “no SQL”

Vs of big data

1. Volume - there is a lot of data
2. Variety - data items look different
3. Velocity - new data is arriving **fast**
4. Veracity - is the data correct?
5. Value - can we do anything with it?

$X \bowtie Y$

Joins



Media

Drawbacks of RDBMS



Flexibility



Scalability

Drawbacks of RDBMS

Volume

- ▶ Maintaining ACID is expensive and not always necessary
- ▶ Sometimes we can deal with minor inconsistencies in our results
- ▶ We also want to be able to partition our data across multiple sites

Variety

- ▶ One single fixed data model makes it harder to incorporate varying data
- ▶ Sometimes when we pull from external sources, we don't know the schema!
- ▶ Changing a schema in a relational database can be expensive

Velocity

- ▶ Storing everything durably to a disk all the time can be prohibitively expensive
- ▶ Sometimes it's ok if we have a low probability of losing data
- ▶ Memory is much cheaper now, and much faster than always going to disk

What is NoSQL?

- ▶ No single accepted definition
- ▶ Flexible schema
(unlike the relational model)
- ▶ Eventual consistency (not ACID)
- ▶ Often (but not always) better at handling **really** big data tasks

Flexible schema

- ▶ Different rows may have different attributes or structure
- ▶ The database often has no understanding of the schema
- ▶ It is up to applications to maintain a consistency in the schema including any denormalization



RDBMS Scalability

- ▶ ACID is hard in a distributed environment
- ▶ Relational databases usually have a primary server and 1+ replicas
- ▶ Writes go to the primary to ensure consistency

ACID vs BASE

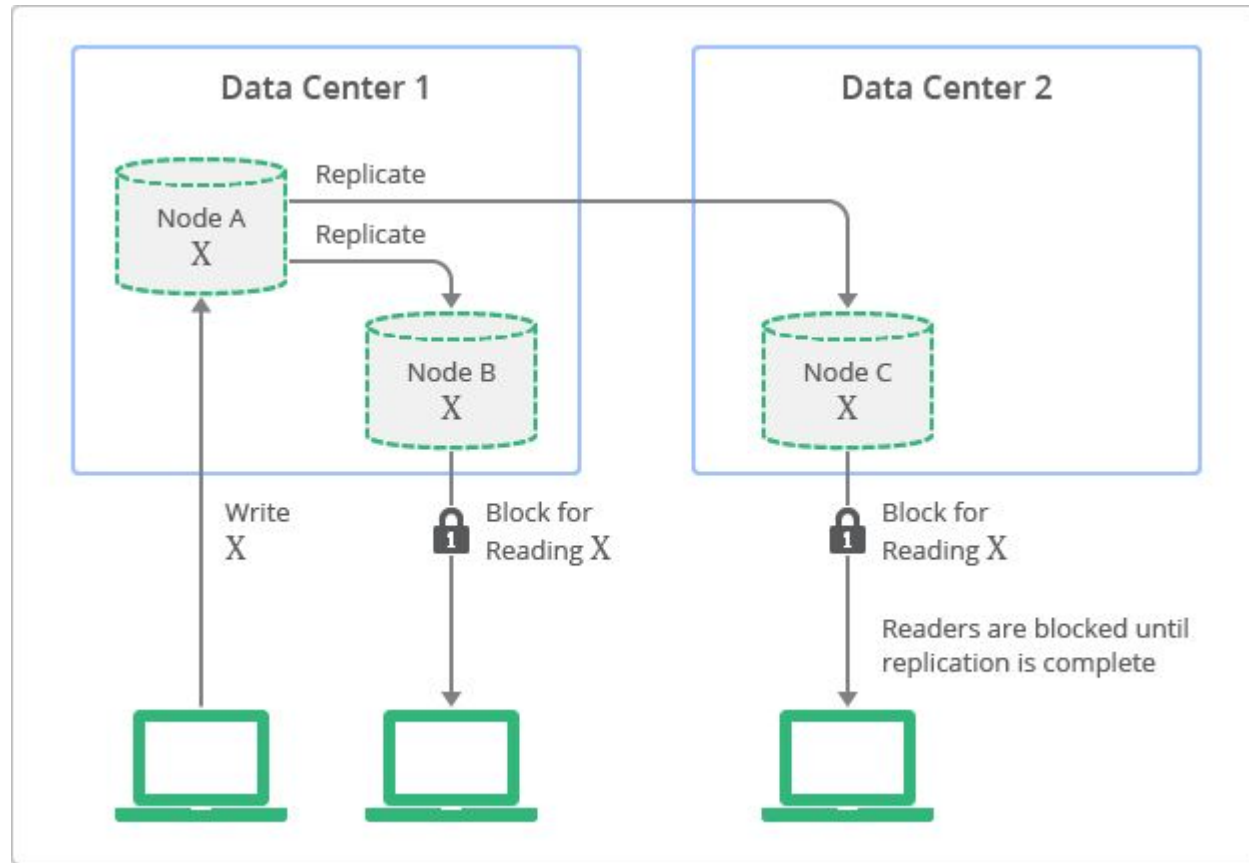
- ▶ ACID databases give us Atomicity, Consistency, Isolation, and Durability
- ▶ This is not universal, but many NoSQL databases choose the **BASE** approach
- ▶ **B**asically **A**vailable
- ▶ **S**oft state
- ▶ **E**ventual consistency

Eventual consistency

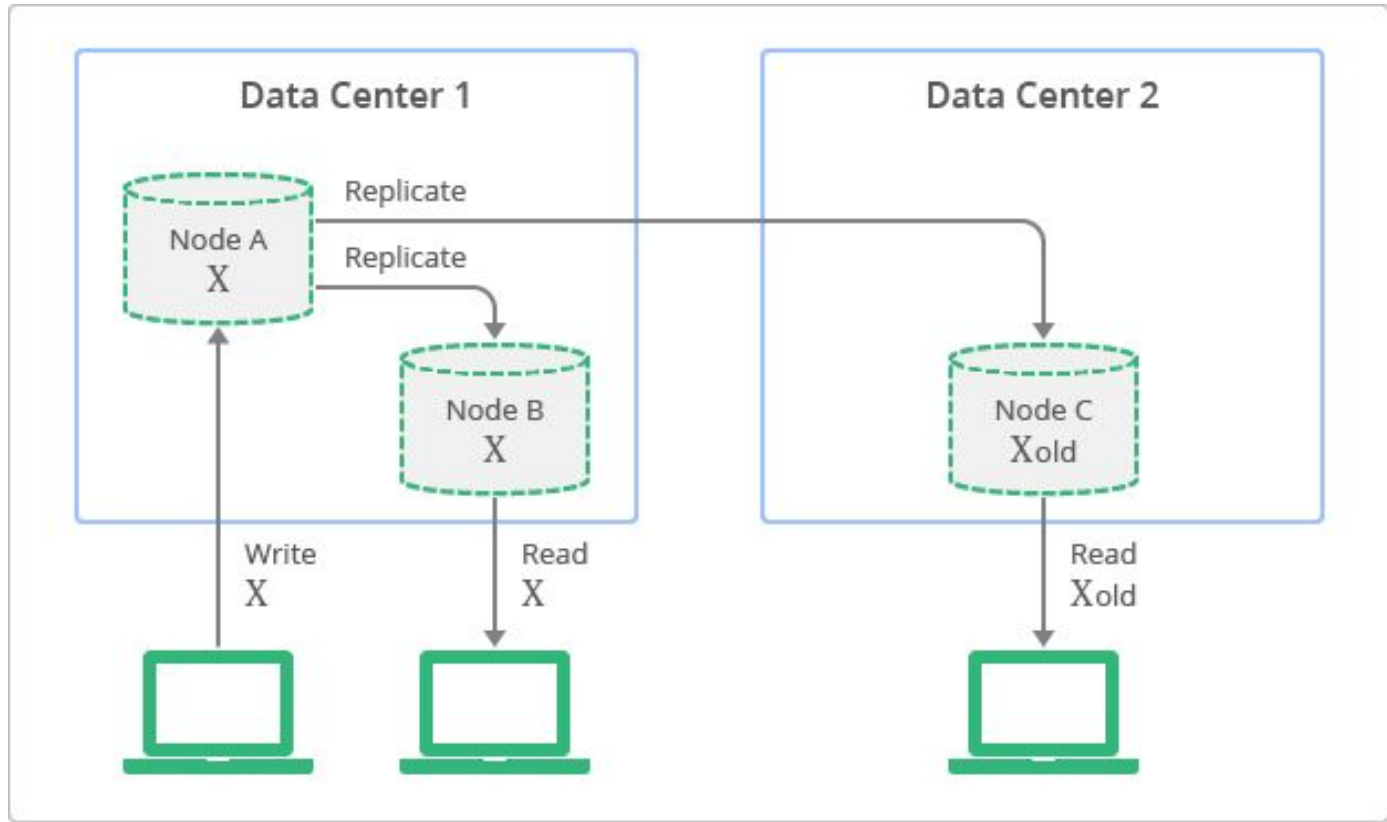
- ▶ Consistency is only guaranteed after some period of time when writes stop
- ▶ This means it is possible that queries will not see the latest data
- ▶ This is commonly implemented by storing data in memory and then lazily sending it to other machines

CAP theorem

- ▶ Consistency, Availability, Partition tolerance, pick two
- ▶ ACID databases are usually CP systems
- ▶ BASE databases are usually AP
- ▶ This distinction is blurry and often systems can be reconfigured to change these tradeoffs



Strong consistency



Eventual consistency

Categories of NoSQL


- ▶ Document stores
- ▶ Key-value databases
- ▶ Wide column stores
- ▶ Graph databases




Sales				
Product	Customer	Date	Sale	
Beer	Thomas	2011-11-25	2	GBP
Beer	Thomas	2011-11-25	2	GBP
Vodka	Thomas	2011-11-25	10	GBP
Whiskey	Christian	2011-11-25	5	GBP
Whiskey	Christian	2011-11-25	5	GBP
Vodka	Alexei	2011-11-25	10	GBP
Vodka	Alexei	2011-11-25	10	GBP

ID	Product
1	Beer
2	Beer
3	Vodka
4	Whiskey
5	Whiskey
6	Vodka
7	Vodka

ID	Customer
1	Thomas
2	Thomas
3	Thomas
4	Christian
5	Christian
6	Alexei
7	Alexei



NoSQL database types




Roadmap

1. Intro to NoSQL
2. Document databases
3. Key-value databases
4. Wide column stores
5. Graph databases



Metadata



Metadata

Metadata



Metadata

Patient		
ssn	firstName	lastName
235-14-7854	Sandra	Smith
192-48-0924	John	Moore
821-13-2108	Laura	Turner

Document databases




mongoDB



FOUNDATIONDB



RethinkDB



MarkLogic®



ArangoDB

Document databases



No standard!



```
{  
  "foo": "bar",  
  "baz": [1, true, 3, "quux"],  
  "corge": {  
    "grault": null  
  }  
}
```

JSON



JSON


- ▶ Taken from JavaScript
- ▶ Contains objects, strings, numbers, arrays, booleans, and null
- ▶ Objects can be arbitrarily nested




mongoDB

MongoDB

Database



Collection



Document

MongoDB hierarchy

```
db.createCollection("bills", [  
  {option: value, ...}])
```

Creating collections

```
{ "firstName": "John",  
  "lastName": "Smith",  
  "dob": new ISODate("1995-10-23"),  
  "address": {  
    "street": "21 2nd St",  
    "city": "New York",  
    "state": "NY",  
    "postalCode": "10021-3100" },  
  "phoneNumbers": [  
    { "type": "home",  
      "number": "212 555-1234" },  
    { "type": "office",  
      "number": "646 555-4567" }  
  ]  
}
```

Sample document



Type	Number	Alias	Notes
Double	1	"double"	
String	2	"string"	
Object	3	"object"	
Array	4	"array"	
Binary data	5	"binData"	
Undefined	6	"undefined"	Deprecated.
ObjectId	7	"objectId"	
Boolean	8	"bool"	
Date	9	"date"	
Null	10	"null"	

BSON

Type	Number	Alias	Notes
Regular Expression	11	"regex"	
DBPointer	12	"dbPointer"	Deprecated.
JavaScript	13	"javascript"	
Symbol	14	"symbol"	Deprecated.
JavaScript (with scope)	15	"javascriptWithScope"	
32-bit integer	16	"int"	
Timestamp	17	"timestamp"	
64-bit integer	18	"long"	
Decimal128	19	"decimal"	New in version 3.4.
Min key	-1	"minKey"	
Max key	127	"maxKey"	

BSON


```
db.bills.insert(  
  { _id : ObjectId("5099803df3f4948bd2f98391"),  
    "address" : { "zipcode" : "14534",  
                  "city"      : "Pittsford",  
                  "state"     : "NY" },  
    "amount" : 756.98,  
    "patient" : "376-97-9845",  
    "id"      : 883,  
    "bDate"   : new ISODate("2010-10-01") } )
```

Need to be
unique in
practice!

Inserting a document

```
db.bills.insert(  
  { "alias" : "BlackJack",  
    "email" : "bj@yeah.com"} )
```



Inserting documents



Structure in practice

```
db.createCollection( "contacts",  
  { validator: { $or:  
    [  
      { phone: { $type: "string" } },  
      { email: { $regex: /@mongodb\.com$/ } },  
      { status: { $in: [ "Unknown", "Incomplete" ] } }  
    ]  
  }  
} )
```

Structure validation

```
db.collection.insertMany(  
  [ <document 1> , <document 2>, ... ],  
  {  
    writeConcern: <document>,  
    ordered: <boolean>  
  }  
)
```

Bulk insert

```
db.bills.find({"patient": "376-97-9845"})  
db.bills.find({"address.zip": "14534"})  
db.bills.find({"amount": {$gte: 750}}) // >=  
db.bills.find({"amount": {$gte: 750}, "address.zipcode": "14534"}) // AND  
db.bills.find({ $or: [{"amount": { $gt: 750 }}, {"address.zip": "14534"} ]})
```

Retrieving data

Comparison

For comparison of different BSON type values, see the [specified BSON comparison order](#).

Name	Description
<code>\$eq</code>	Matches values that are equal to a specified value.
<code>\$gt</code>	Matches values that are greater than a specified value.
<code>\$gte</code>	Matches values that are greater than or equal to a specified value.
<code>\$in</code>	Matches any of the values specified in an array.
<code>\$lt</code>	Matches values that are less than a specified value.
<code>\$lte</code>	Matches values that are less than or equal to a specified value.
<code>\$ne</code>	Matches all values that are not equal to a specified value.
<code>\$nin</code>	Matches none of the values specified in an array.

Query operators

Logical

Name	Description
<code>\$and</code>	Joins query clauses with a logical AND returns all documents that match the conditions of both clauses.
<code>\$not</code>	Inverts the effect of a query expression and returns documents that do <i>not</i> match the query expression.
<code>\$nor</code>	Joins query clauses with a logical NOR returns all documents that fail to match both clauses.
<code>\$or</code>	Joins query clauses with a logical OR returns all documents that match the conditions of either clause.

Query operators

Element

Name	Description
<code>\$exists</code>	Matches documents that have the specified field.
<code>\$type</code>	Selects documents if a field is of the specified type.

Evaluation

Name	Description
<code>\$mod</code>	Performs a modulo operation on the value of a field and selects documents with a specified result.
<code>\$regex</code>	Selects documents where values match a specified regular expression.
<code>\$text</code>	Performs text search.
<code>\$where</code>	Matches documents that satisfy a JavaScript expression.

Query operators

Array

Name	Description
<code>\$all</code>	Matches arrays that contain all elements specified in the query.
<code>\$elemMatch</code>	Selects documents if element in the array field matches all the specified <code>\$elemMatch</code> conditions.
<code>\$size</code>	Selects documents if the array field is a specified size.

Bitwise

Name	Description
<code>\$bitsAllClear</code>	Matches numeric or binary values in which a set of bit positions <i>all</i> have a value of <code>0</code> .
<code>\$bitsAllSet</code>	Matches numeric or binary values in which a set of bit positions <i>all</i> have a value of <code>1</code> .
<code>\$bitsAnyClear</code>	Matches numeric or binary values in which <i>any</i> bit from a set of bit positions has a value of <code>0</code> .
<code>\$bitsAnySet</code>	Matches numeric or binary values in which <i>any</i> bit from a set of bit positions has a value of <code>1</code> .

Query operators

MongoDB does NOT support SQL, however I will include some SQL statements that are roughly equivalent. When asked to write a query in MongoDB, **always** use the MongoDB query language!

Using SQL will result in **zero** marks.



MongoDB vs SQL

```
db.bills.find( {"patient": "376-97-9845"},  
               {"patient": 1, "id": 1 } )
```

If this were SQL (NOT supported by MongoDB)...

```
SELECT patient, id FROM bills WHERE  
       patient="376-97-9845"
```

Projection

```
db.bills.find( {"patient": "376-97-9845"},  
               { "patient": 1, "id": 1 } )
```

A diagram consisting of three red arrows. One arrow points from the string "376-97-9845" in the MongoDB query to the string "376-97-9845" in the SQL query. A second arrow points from the "patient" field in the MongoDB query's second argument to the "patient" field in the SQL query's SELECT clause. A third arrow points from the "id" field in the MongoDB query's second argument to the "id" field in the SQL query's SELECT clause.

If this were SQL (NOT supported by MongoDB)...

```
SELECT patient, id FROM bills WHERE  
patient='376-97-9845'
```

Projection

```
db.bills.find().sort(  
    {"patient": 1,  
     "address.zip": -1})
```

```
SELECT * FROM bills ORDER by  
    patient ASC, address.zip DESC
```

Sorting

```
db.bills.update(  
    {"patient": "376-97-9845"},  
    {$set: {"address.zip": "14626",  
            "address.city": "Greece"}})
```

```
UPDATE bills SET address.zip="14626",  
address.city="Greece" WHERE  
patient="376-97-9845"
```

Updating documents

```
db.bills.remove(  
    {"patient": "376-97-9845"})
```

```
DELETE FROM WHERE patient="376-97-9845"
```

Removing documents


```
{field: {$exists: <boolean>}}
```

Not the same as **EXISTS** in SQL!

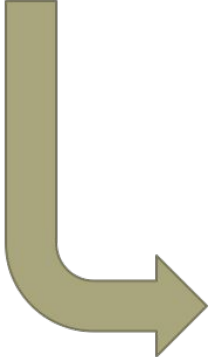
Exists

```
{ $project: { <specification(s)> } }
```

- **<field>: <1 or true>, <field>:<0 or false>**
 - Inclusion/exclusion of a field
- **_id: <0 or false>**
 - Suppression of the _id field
- **<field>: <expression>**
 - Add a new field or reset the value of an existing field.
- If you specify the exclusion of a field other than _id, you cannot employ any other \$project specification forms.

Projection stage

```
db.bills.aggregate( [  
  { $project : { "address.zipcode" : 1,  
    "amount" : 1,  
    "patient" : 1 } } ] )
```



```
{ _id : "5099803df3f4948bd2f98391",  
  "address" : { "zipcode" : "14534" },  
  "amount" : 756.98,  
  "patient" : "376-97-9845" }
```

Projection stage

```
db.bills.aggregate( [  
  { $project : { "address.zipcode" : 1,  
    "amount" : 1,  
    "patient" : 1,  
    _id : 0 } } ] )
```

Documents
with no _ids!

Projection stage

```
{ $match: { <query> } }
```

Match stage

```
db.bills.aggregate([  
  {$match : {"address.zip" : "14534"}}])
```

same as

```
db.bills.find({"address.zip" : "14534"})
```

Match stage

```
{ $group: { _id: <expression>, <field1>: { <accumulator1> : <expression1> }, ... } }
```

Collection
↓
db.orders.aggregate([
 \$match stage → { \$match: { status: "A" } },
 \$group stage → { \$group: { _id: "\$cust_id", total: { \$sum: "\$amount" } } }
])

{ cust_id: "A123", amount: 500, status: "A" }
{ cust_id: "A123", amount: 250, status: "A" }
{ cust_id: "B212", amount: 200, status: "A" }
{ cust_id: "A123", amount: 300, status: "D" }

orders

\$match →

{ cust_id: "A123", amount: 500, status: "A" }
{ cust_id: "A123", amount: 250, status: "A" }
{ cust_id: "B212", amount: 200, status: "A" }

\$group →

Results
{ _id: "A123", total: 750 }
{ _id: "B212", total: 200 }

SELECT cust_id AS _id,
SUM(amount) FROM ORDERS
WHERE status="A"
GROUP BY cust_id

Example

\$sum	Sum
\$avg	Average
\$first	First value in group
\$last	Last value in group
\$max	Maximum
\$min	Minimum
\$push	Array of all values in group
\$addToSet	Distinct array of group values
\$stdDevPop	Population standard deviation
\$stdDevSamp	Sample standard deviation

Accumulators

```
db.bills.aggregate(  
  [$group: {_id : null,  
            avgAmount: {$avg: "$amount"}}])
```

```
SELECT AVG(amount) FROM bills
```

Grouping by null

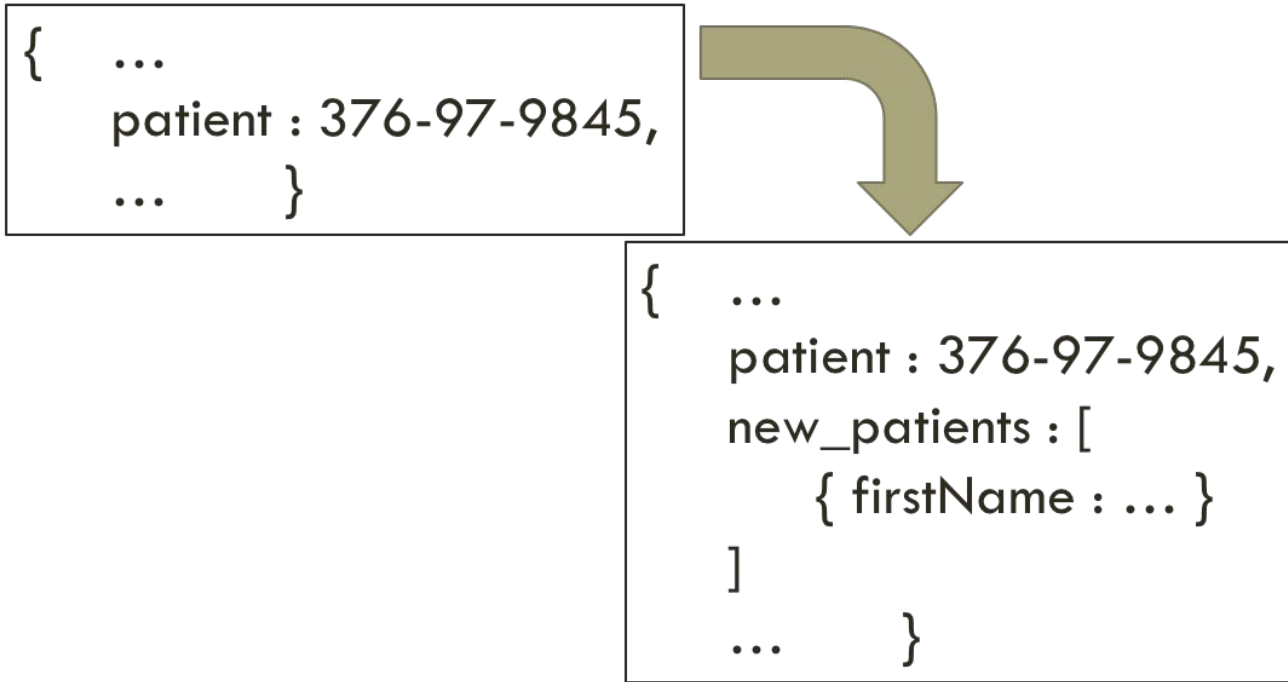
```
{
  $lookup:
  {
    from: <collection to join>,
    localField: <field from the input documents>,
    foreignField: <field from the documents of the "from" collection>,
    as: <output array field>
  }
}
```

Lookup stage

```
db.bills.aggregate(  
    [$lookup: {from: patients,  
                localField : "patient",  
                foreignField : "_id",  
                as: "new_patients"}])
```

```
SELECT * FROM bills JOIN patients AS new_patients  
ON bills.patient = new_patients._id
```

Lookup stage



Lookup stage

```
db.patients.insert(  
  { _id : "376-97-9845",  
    "firstName" : "Jennifer",  
    "visits" : [758, 345],  
    ... } )
```

```
db.visits.insert(  
  { _id : 345,  
    "doctor" : "893-12-8934",  
    "otherNotes" : "fever",  
    ... } )  
db.visits.insert(  
  { _id : 758,  
    "doctor" : "9094-56-9292",  
    "otherNotes" : "neck",  
    ... } )
```

Lookup stage

```
db.patients.aggregate(  
    [$lookup: {from: visits,  
                localField: "visits",  
                foreignField: "_id",  
                as: "visitsInfo"}])
```

```
SELECT * FROM patients JOIN visits AS  
visitsInfo ON patients.visits = visitsInfo._id
```

Lookup stage

```
{ _id : "376-97-9845",  
  "firstName" : "Jennifer",  
  "visits" : [758, 345],  
  "visitsInfo" : [  
    { _id : 345,  
      "doctor" : "893-12-8934",  
      "otherNotes" : "fever",  
      ... },
```

```
{ _id : 758,  
  "doctor" : "9094-56-9292",  
  "otherNotes" : "neck",  
  ... }  
],  
... }
```

Lookup stage


```
{ $unwind: <field path> }
```

Unwind stage

Consider an `inventory` with the following document:

```
{ "_id" : 1, "item" : "ABC1", sizes: [ "S", "M", "L" ] }
```

The following aggregation uses the `$unwind` stage to output a document for each element in the `sizes` array:

```
db.inventory.aggregate( [ { $unwind : "$sizes" } ] )
```

The operation returns the following results:

```
{ "_id" : 1, "item" : "ABC1", "sizes" : "S" }  
{ "_id" : 1, "item" : "ABC1", "sizes" : "M" }  
{ "_id" : 1, "item" : "ABC1", "sizes" : "L" }
```

Documents
with the
same `_ids`!

Unwind stage

```
db.records.createIndex( { score: 1 } )
```

Indexes

Roadmap

1. Intro to NoSQL
2. Document databases
3. **Key-value databases**
4. Wide column stores
5. Graph databases



Key-value (KV) databases



No standard!

Key-value databases

- ▶ Updates to the value for a single key are usually atomic
- ▶ Many KV databases allow for transactions which use multiple keys
- ▶ Values have limited structure

Upsides

- ▶ Key-value databases are generally easier to run in a distributed fashion
- ▶ Queries and updates usually very fast
- ▶ Any type of data in any structure can be stored as a value

Downsides



- ▶ **Very** simple queries (usually just get a value given a key, sometimes a range)
- ▶ No referential integrity
- ▶ Limited transactional capabilities
- ▶ No schema to understand the data




Redis



Replicas



Primary



Scalability

Redis data model

- ▶ Basically a huge distributed hash table (DHT) with little structure to values
- ▶ All values are identified by a key which is a simple string
- ▶ If we want more structure in our keys, it has to be defined by our application e.g., user 3 could have the key “user:3”



Redis values

- ▶ Commonly in key-value stores, values are just an arbitrary blob of data
- ▶ Redis (and some other KV stores) allows some structure:
 - ▷ Lists
 - ▷ Sets
 - ▷ Hashes



```
SET foo bar
```

```
SADD news:1000:tags 1 2 5 77
```

```
LPUSH mylist first
```

```
HMSET user:1 username antirez birthyear 1977
```

Inserting data in Redis



```
GET foo => bar
```

```
SMEMBERS news:1000:tags => 1 2 5 77
```

```
LPOP mylist => first
```

```
HMGET user:1 username birthyear =>  
antirez 1977
```

Retrieving data in Redis

When to use KV databases

- ▶ When you need something *really* fast
- ▶ When your data does not have a lot of structure/relationships
- ▶ For simple caching of data which is pulled from another source

Roadmap


1. Intro to NoSQL
2. Document databases
3. Key-value databases
4. **Wide column stores**
5. Graph databases

Patient		
ssn	firstName	lastName
235-14-7854	Sandra	Smith
192-48-0924	John	Moore
821-13-2108	Laura	Turner

Wide-column stores

Patient		
ssn	firstName	lastName
A235-14-7854	B Sandra	C Smith
D192-48-0924	E John	F Moore
G821-13-2108	H Laura	I Turner

SELECT ssn
FROM Patient;



Wide-column stores

"follows" column family

	Follows			
Row Key	gwasington	jadams	tjefferson	wmckinley
gwasington		1		
jadams	1		1	
tjefferson	1	1		1
wmckinley			1	

Multiple versions

Wide-column stores

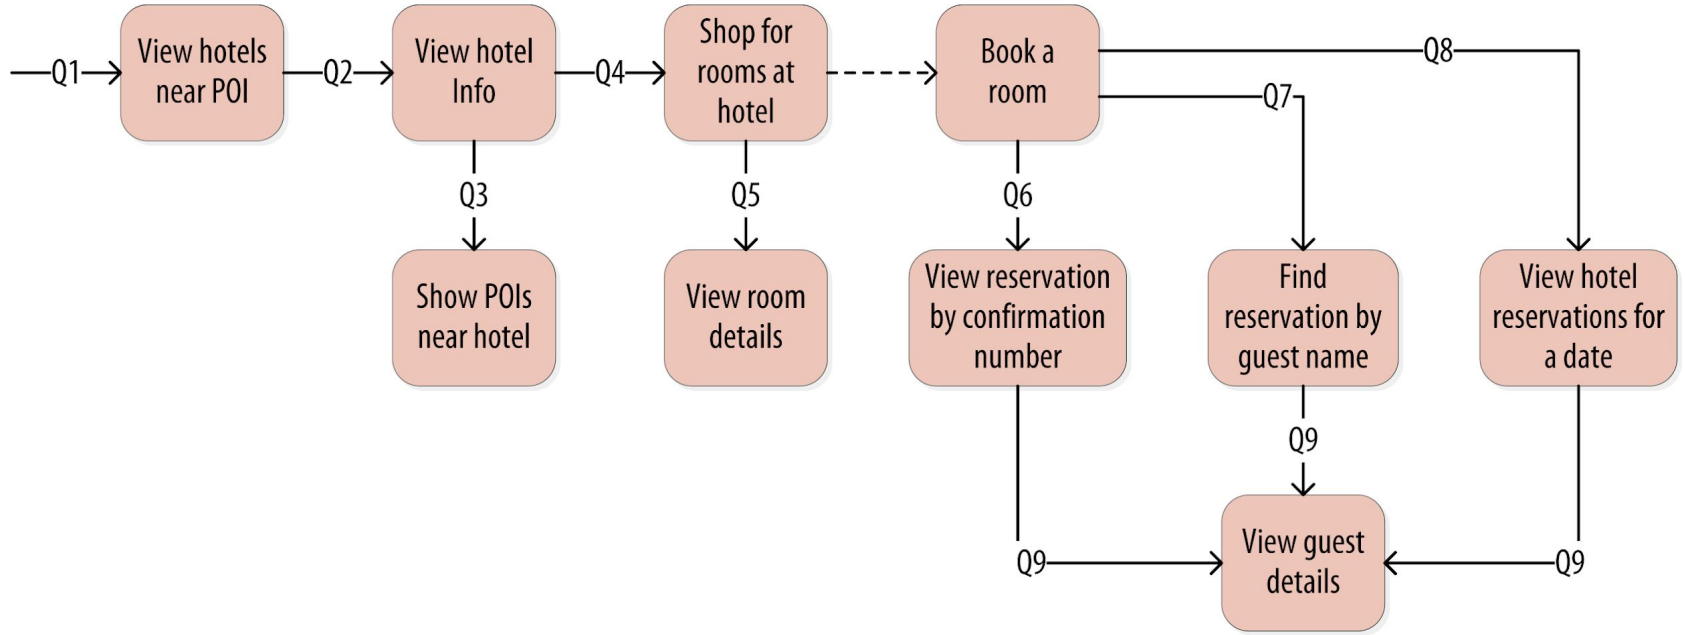


No standard!

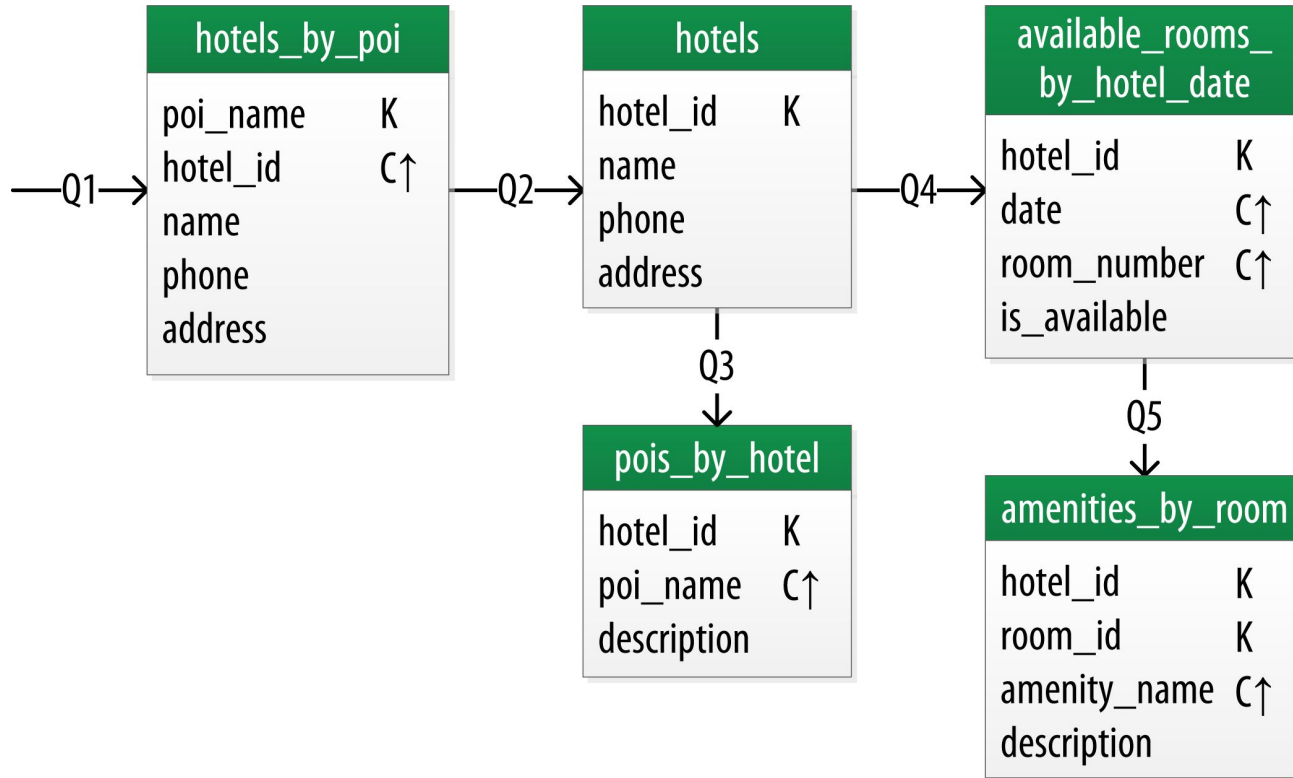


cassandra

Cassandra



Logical model



Logical model



```
SELECT name, phone, address  
FROM hotels  
WHERE  
    hotel_id = "127" OR  
    hotel_id = "349";
```

Querying using CQL

```
SELECT name, phone, address  
FROM hotels  
WHERE  
    phone = "12345678910";
```



Querying using CQL



Scalability

Consistent hashing

- ▶ To find what data a node should hold, hash values are organized in a ring
- ▶ Each node is responsible for a portion of the ring
- ▶ When nodes are added or removed, this minimizes data movement

Upsides

- ▶ Very fast for writes
- ▶ Some structure to data and queries are somewhat familiar
- ▶ Highly available since there is no single point of failure

Downsides

- ▶ Limited queries (no joins!)
- ▶ No referential integrity
- ▶ Eventual consistency

Denormalization

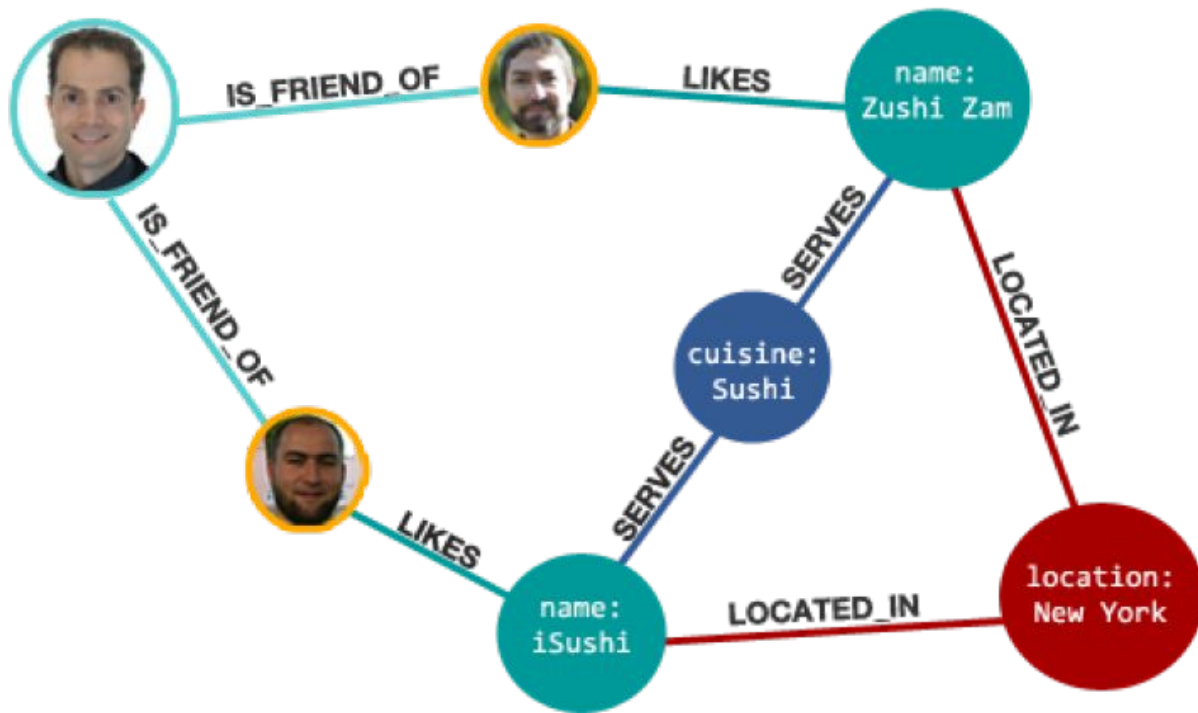
- ▶ Since the queries that can be issued are very simple, denormalization is often required
- ▶ A common approach is to create a column family for each type of query and pre-compute results
- ▶ Some systems support materialized views and indexes

When to use wide column stores


- ▶ Especially great for time series data (e.g. sensor measurements, logs)
- ▶ Simple data with predictable structure
- ▶ When you are writing a lot of new data but rarely updating old data
- ▶ If you don't need transactions or complex queries

Roadmap

1. Intro to NoSQL
2. Document databases
3. Key-value databases
4. Wide column stores
5. Graph databases



Graph databases



Unbounded queries

Graph databases

Friend	
person	contact
John	Matthew
Matthew	Jennifer
Jennifer	Paul

```
SELECT person, contact  
FROM Friend
```


```
WITH  
SELECT f1.person, f2.contact  
FROM Friend AS f1 JOIN  
      Friend AS f2 ON  
      f1.contact = f2.person  
WHERE  
      f1.person = "John" AND  
      f2.contact = "Paul";
```

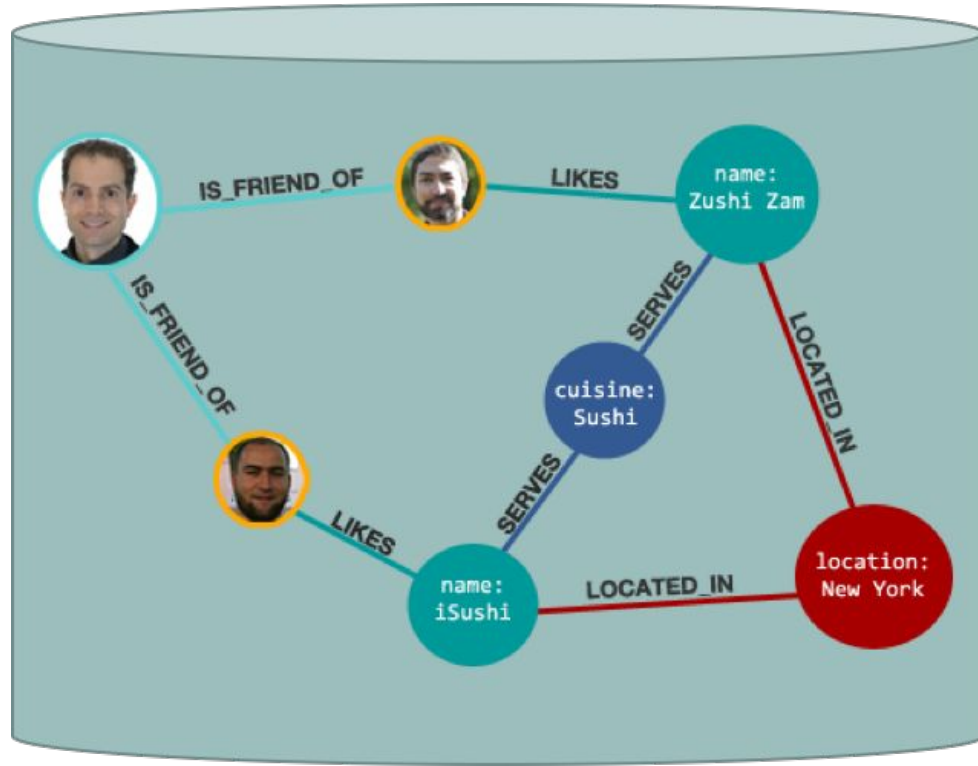
Unbounded SQL queries

Friend	
person	contact
John	Matthew
Matthew	Jennifer
Jennifer	Paul

```
WITH RECURSIVE friends(name)
AS (VALUES('John') UNION
SELECT contact from friend JOIN
friends ON person=name)
SELECT * from friends OFFSET 1;
```

Unbounded SQL queries





Databases and graphs

PERSON, PATIENT

[984148583247]

ssn: 235-14-7854

first: Sandra

last: Smith

age: 38

Nodes


```
CREATE (n:PERSON {first: 'Sandra'})
```

```
Map<String, Object> attributesSandra =  
    new HashMap<>();  
attributesSandra.put("first", "Sandra");  
long sandraNode = inserter.createNode(  
    attributesSandra, Label.label("PERSON"));
```

Creating a node

```
inserter.createNode(1L,  
    attributesSandra, Label.label("PERSON"),  
    Label.label("PATIENT"));
```

Internal node IDs

PERSON, PATIENT

ssn: 235-14-7854

first: Sandra

last: Smith

PERSON, DOCTOR

ssn: 621-11-0923

first: John

last: Moore

PRIMARY

start: 2015-08-23

Relationships

```
CREATE (p) [:PRIMARY {start: '2015-08-23'}] -> (d)
```

```
Map<String, Object> att =
```

```
    new HashMap<>();
```

```
att.put("start", "2015-08-23");
```

```
inserter.createRelationship(sandra, john,
```

```
    RelationshipType.withName("PRIMARY"), att);
```

Creating edges

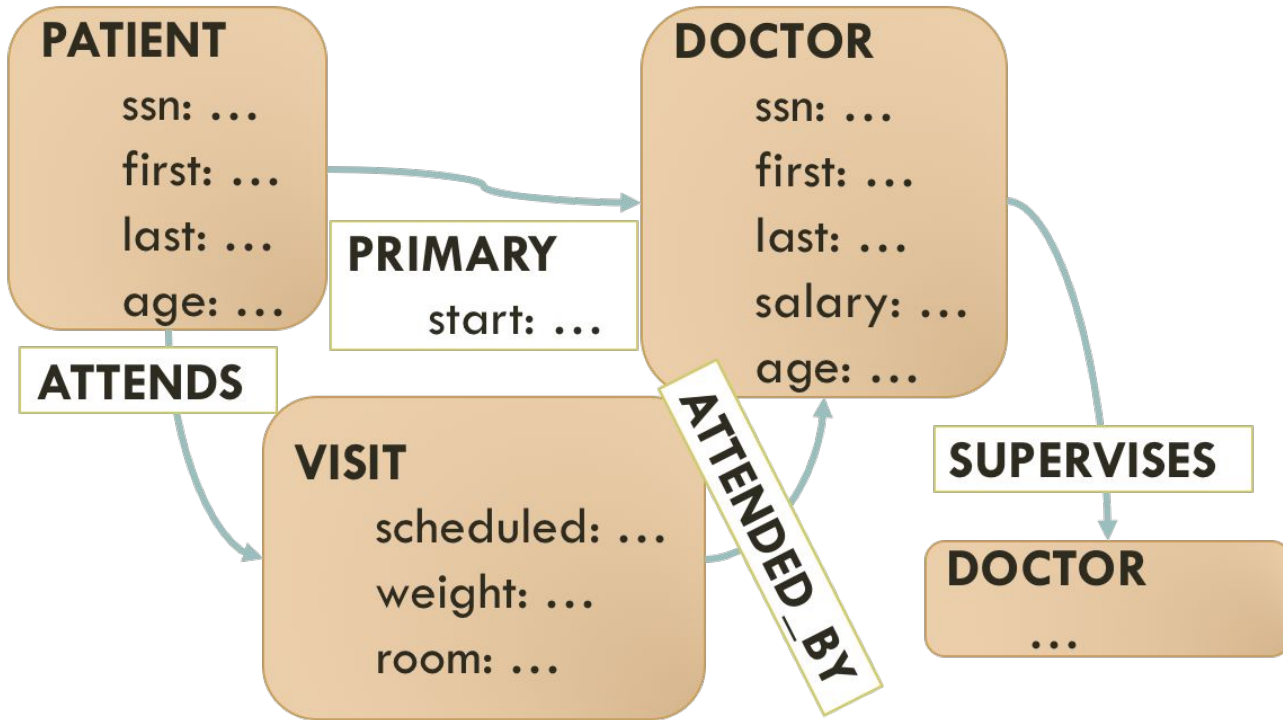


Logical model

```
CREATE CONSTRAINT ON (p:PATIENT)  
  ASSERT p.ssn IS UNIQUE;
```

```
CREATE CONSTRAINT ON (p:PATIENT)  
  ASSERT EXISTS(p.ssn);
```

Logical model



Sample data

```
MATCH (p:PATIENT)
WHERE p.age > 30
RETURN p.last
ORDER BY p.last ASC
```

Single node queries


```
MATCH (p:PATIENT) --> (d:Doctor)
WHERE p.age > 30 AND
      d.salary > $100K
RETURN (p)
```

Queries with edges

```
MATCH (p:PATIENT) -[:PRIMARY]- (d:Doctor)
WHERE p.age > 30 AND
      d.salary > $100K
RETURN (p)
```

Queries with edges

```
MATCH (d:DOCTOR)  
WHERE d.age > 30  
RETURN AVG(d.salary)
```

Aggregation