

CSCI-620

Data Integration

Data integration

- ▶ We often need to combine data from multiple sources
- ▶ There are multiple strategies for this
 - ▷ Data warehousing
 - ▷ Federation
 - ▷ Data lakes

Data Warehouse

- ▶ Where data is stored in a form suitable for analysis and reporting
- ▶ Data is often denormalized to make reports faster to generate
- ▶ Warehouses usually contain historical data (possibly from multiple sources) that is refreshed periodically

Extract Transform Load

- ▶ **Extract** data from multiple sources
useful for reporting
- ▶ **Transform** the data into a format more
suitable for analysis
- ▶ **Load** the data into the warehouse

Views

- ▶ Creates shortcuts for repeated or common queries
- ▶ Views may be *materialized* (stored) to make complex queries faster
- ▶ Materialized views must be manually updated or they will return old data
- ▶ Some systems will automatically rewrite queries to use materialized views

```
CREATE VIEW doctors_with_supervisor  
AS SELECT d.*, supervisor, s.firstName ||  
' ' || s.lastName AS supervisorName  
FROM Doctor d JOIN SupervisedBy ON  
d.ssn=supervisee JOIN Doctor s ON  
s.ssn=supervisor;
```

CREATE VIEW

```
SELECT s.* FROM salaries_by_dept s JOIN  
Department d ON d.id=s.department;
```

```
SELECT s.* FROM (SELECT Doctor.department,  
MIN(Doctor.salary), MAX(Doctor.salary) FROM  
Doctor GROUP BY Doctor.department) s JOIN  
Department d ON d.id=s.department;
```

View expansion

```
EXPLAIN SELECT s.* FROM salaries_by_dept s JOIN  
Department d ON d.id=s.department;
```

Hash Join

Hash Cond: (d.id = doctor.department)

-> Seq Scan on department d

-> Hash

-> HashAggregate

Group Key: doctor.department

-> Seq Scan on doctor

View expansion


```
CREATE MATERIALIZED VIEW  
salaries_by_dept AS SELECT  
department, MIN(salary),  
MAX(salary) FROM Doctor GROUP BY  
department;
```

CREATE MATERIALIZED VIEW

```
EXPLAIN SELECT s.* FROM salaries_by_dept s JOIN  
Department d ON d.id=s.department;
```

Hash Join

Hash Cond: (s.department = d.id)

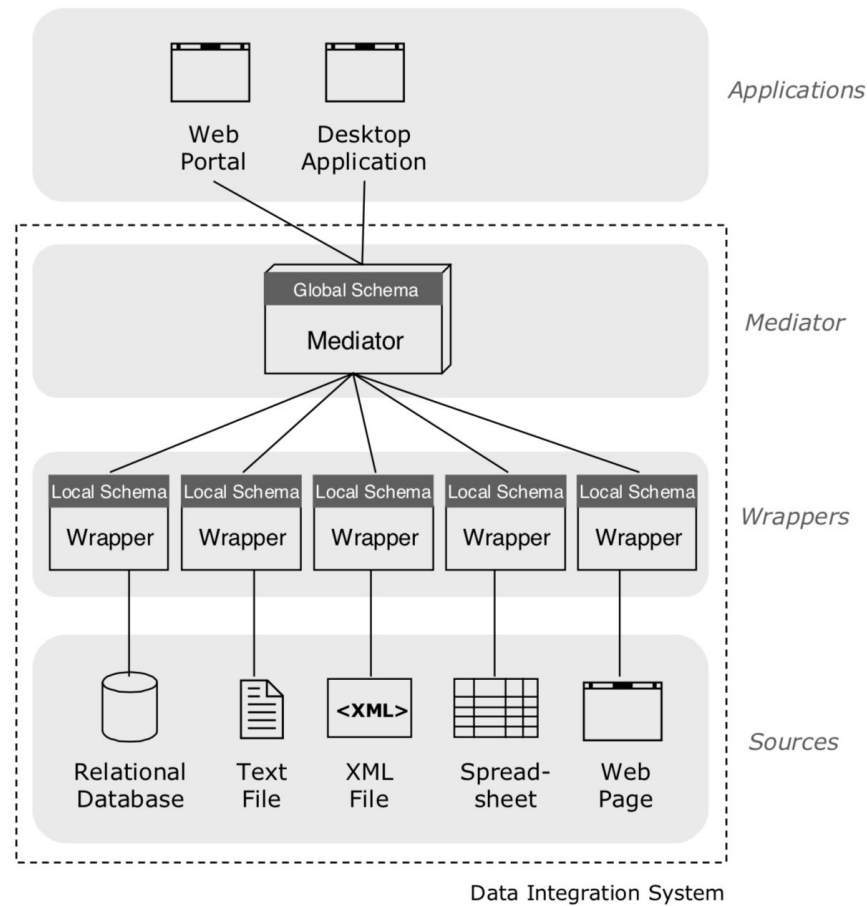
-> Seq Scan on salaries_by_dept s

-> Hash

-> Seq Scan on department d

Global schema

- ▶ When integrating data, our goal is to produce a *global* or shared schema
- ▶ We can write the global schema or each of our local schemas as a view



Data Integration System

View-based data integration



L_{1a}: Prentice Hall

□ PHBook

- ISBN
- title
- authorID
- sug_retail
- format

□ PHAuthor

- authorID
- authorName

L_{1b}: Prentice Hall_{condensed}

□ PHBook_{condensed}

- ISBN
- title
- authorName
- sug_retail

L₂: Barnes & Noble

○ BNNewDeliveries

- ISBN
- title
- instock

G: Book Portal

△ Book

- ISBN
- title
- sug_retail
- author
- publisher

△ Book_Price

- ISBN
- seller
- price
- instock

(a) Local Schemas

(b) Global Schema

Local and Global Schemas

Example

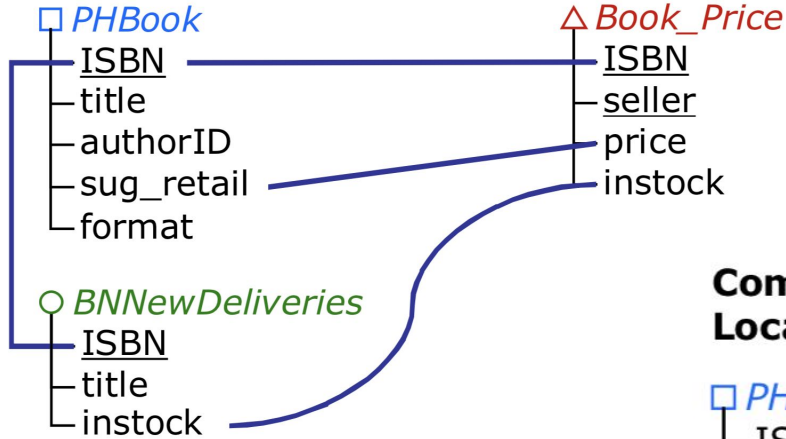
Global as view (GAV)

- ▶ If sources are unlikely to change, we can express our global schema G as a view over each data source L_i
- ▶ Queries on the global schema are easy to execute based on this mapping
- ▶ Adding a new source later is hard




Combined Local Schemas

Global Schema Relation



Combined Local Schemas

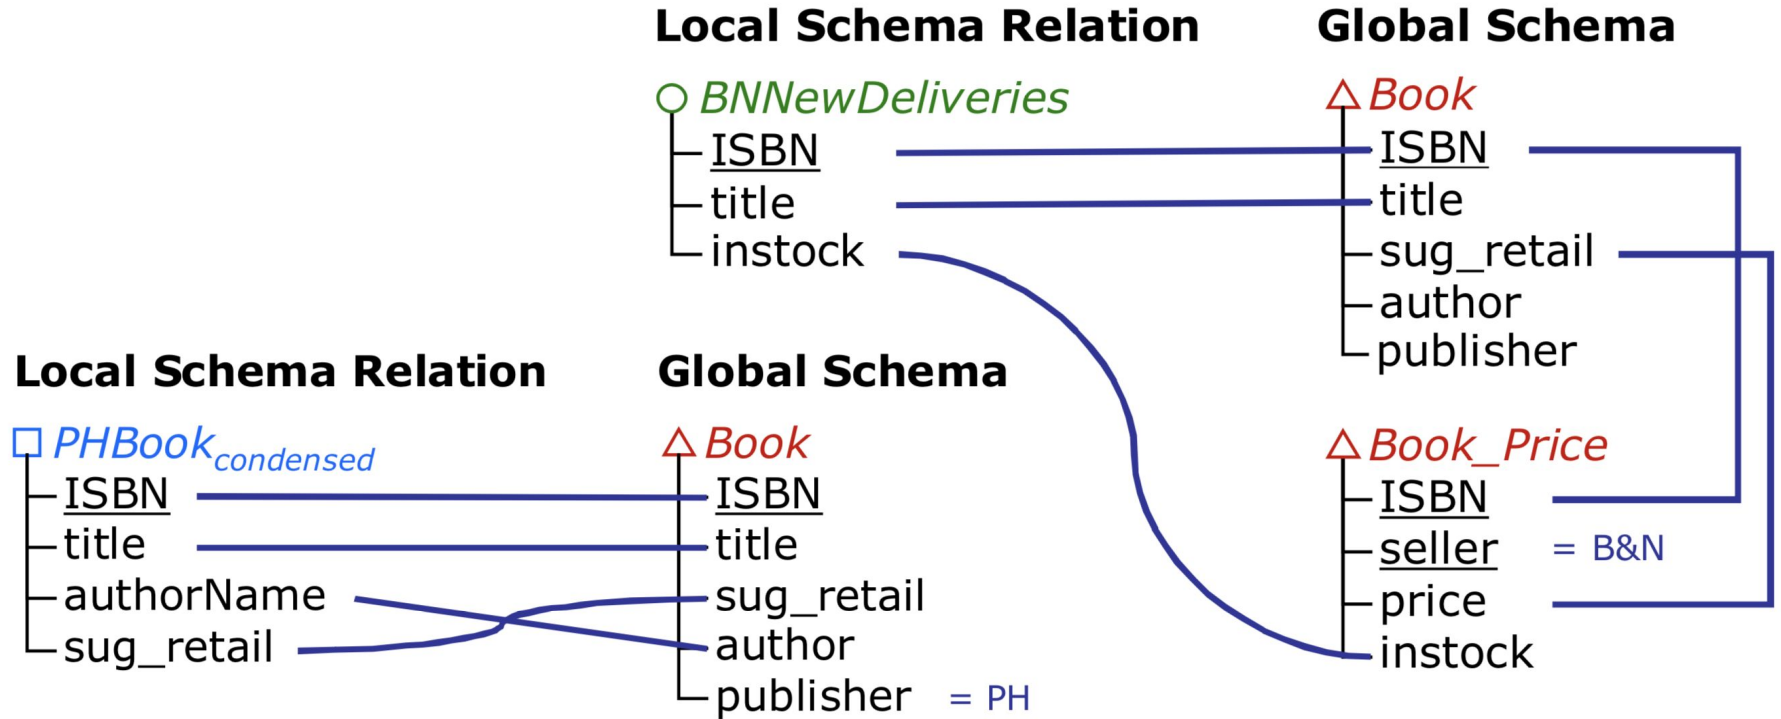
Global Schema Relation



Example

Local as view (LAV)

- ▶ When sources change frequently, we can instead express each source as a view over the global schema
- ▶ Processing queries is harder
- ▶ Adding new sources is easy since we only need to write a single view



Example

**Global**

Cust (ID, firstName, lastName, ...)
CustPhones(ID, Type, PhoneNum, ...)

Source 1

Customers (ID, firstName, lastName, homePhone, cellPhone, ...)

Source 2

Customers (ID, firstName, lastName, ...)
CustomersPhones(ID, Type, PhoneNum)

Another Example

Extract Transform Load

- ▶ **Extract** data from multiple sources
useful for reporting
- ▶ **Transform** the data into a format more
suitable for analysis
- ▶ **Load** the data into the warehouse



Federation

- ▶ A *federated* database system allows one database to connect to other systems
- ▶ For example, PostgreSQL has foreign data wrappers for this purpose
- ▶ Queries can then join with data on other systems with each executing part of the query

```
CREATE FOREIGN TABLE employees (  
    id integer,  
    name text,  
    address text)  
SERVER mysql_svr  
OPTIONS (table 'hr.employees');
```


Foreign data wrapper

Data lakes

- ▶ One problem with data warehousing is that ETL can be time-consuming
- ▶ If we have a lot of data, then it may not be possible to keep up
- ▶ Instead, a data lake just collects large amounts of data without any transformation step

Data lakes

- ▶ Skipping the transform steps solves the velocity problem
- ▶ However, data is now in many different formats making queries challenging
- ▶ We need to make sure we maintain appropriate *metadata* for the data



Data lake

Other integration problems

- ▶ **Schema matching**
Aligning schemas from different sources
- ▶ **Entity resolution**
Identifying the same real world thing across sources



Schema matching


- ▶ Data may come from different systems but we want to connect them together
- ▶ This has several challenges
 - ▷ Syntax - languages used by the DBs
 - ▷ Structure - layout of records
 - ▷ Semantics - description used

```
CREATE TABLE users (  
    id INT PRIMARY KEY,  
    username VARCHAR(50));
```

```
db.createCollection("users", {  
    validator: {$jsonSchema: {  
        bsonType: "object",  
        required: ["username"],  
        properties: {username: {bsonType: "string"}}  
    }})
```

Syntactic heterogeneity


Customer			
id	street	city	zip
3	One Lomb Memorial Drive	Rochester	14623



Customer	
id	address
3	One Lomb Memorial Drive, Rochester, 14623

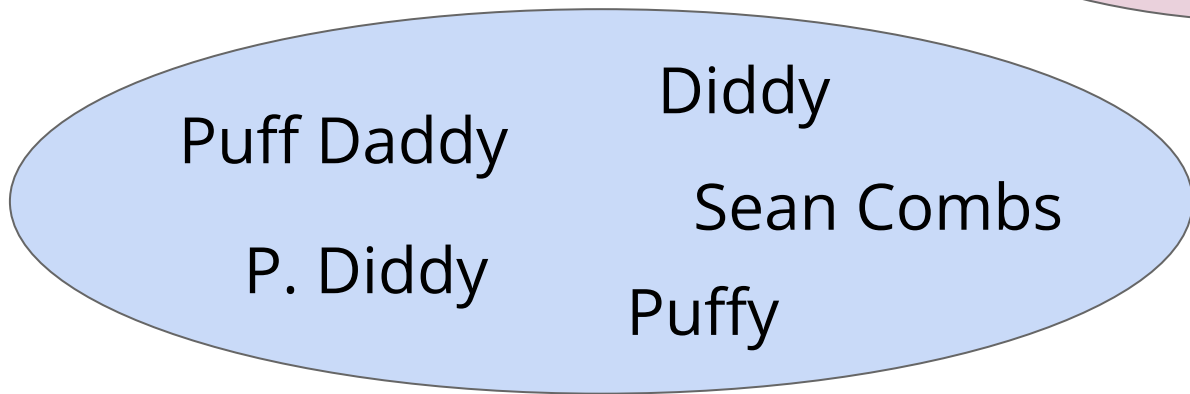

Structural heterogeneity

Customer			
id	company	balance	username
3	RIT	37000	rit



Client			
account	business	amount	login
3	RIT	37000	rit

Semantic heterogeneity



Entity Resolution

Ontologies

- ▶ Ontologies provide a standardized representation of semantics that can be shared across data sources
- ▶ Several large scale instances exist:
 - ▷ Google Knowledge Graph
 - ▷ DBpedia
 - ▷ Schema.org